

10.2. Hash Tables 415

Hash Codes in Python

The standard mechanism for computing hash codes in Python is a built-in function with signature `hash(x)` that returns an integer value that serves as the hash code for object `x`. However, only immutable data types are deemed hashable in Python. This restriction is meant to ensure that a particular object's hash code remains constant during that object's lifespan. This is an important property for an object's use as a key in a hash table. A problem could occur if a key were inserted into the hash table, yet a later search were performed. Among Python's built-in data types, the immutable `int`, `float`, `str`, `tuple`, and `frozenset` classes produce robust hash codes, via the `hash` function, using techniques similar to those discussed earlier in this section. Hash codes for character strings are well crafted based on the order of the elements. Hash codes for tuples are computed with a similar technique based upon a combination of the hash codes of the individual elements of the tuple. When hashing a `frozenset`, the order of the elements should be irrelevant, and so a natural option is to compute the exclusive-or of the individual hash codes. For an instance `x` of a mutable type, such as a list, a `TypeError` is raised. Instances of user-defined classes are treated as unhashable by default, with a `TypeError` raised by the `hash` function. However, a function that computes hash codes can be implemented in the form of a special method named

`hash`

within

a class. The returned hash code should reflect the immutable attributes of an instance. It is common to return a hash code composed of three numeric red, green, and blue components might implement the method as:

```
def
```

```
hash
```

```
(self):
```

```
    return hash( ( self.
```

```
red, self.
```

```
green, self.
```

```
blue) ) # hash combined tuple
```

An important rule to obey is that if a class defines equivalence through

```
eq
```

```
,
```

then any implementation of

```
hash
```

must be consistent, in that if `x == y`, then

`hash(x) == hash(y)`. This is important because if two instances are considered

to be equivalent and one is used as a key in a hash table, a search for the second instance should result in the dictionary. true, it ensures that `hash(5)` and `hash(5.0)` are the same.

Hash Codes

The first action that a hash function performs is to take an arbitrary key k in our map and compute an integer that is called the hash code for k ; this integer need not be in the range $[0, N-1]$, and may even be negative. We desire that the set of hash codes assigned to our keys should avoid collisions as much as possible. For if the hash codes of our keys cause collisions, then there is no hope for our compression function to avoid them. In this subsection, we begin by discussing the theory of hash codes. Following that, we discuss treating the bit representation as an integer.

To begin, we note that, for any data type X that is represented using at most as many bits as our integer hash codes, we can simply take as a hash code for x an integer interpretation of its bits. For example, the hash code for key 314 could simply be 314. The hash code for a floating-point number such as 3.14 could be based upon an interpretation of the bits of the floating-point representation as an integer.

For a type whose bit representation is longer than a desired hash code, the above scheme is not immediately applicable. For example, Python relies on 32-bit hash codes. If a floating-point number uses a 64-bit representation, its bits cannot be viewed directly as a hash code. One possibility is to use only the high-order 32 bits (or the low-order 32 bits). This hash code, of course, ignores half of the information present in the original key, and is thus not ideal.

A better approach is to combine in some way the high-order and low-order portions of a 64-bit key to form a 32-bit hash code, which takes all the original bits into consideration. A simple implementation is to view the binary representation as an n -tuple $(x_0, x_1, \dots, x_{n-1})$ of 32-bit integers, for example, by forming a hash code for x as

$$h(x) = x_0 \oplus x_1 \oplus \dots \oplus x_{n-1},$$

where the \oplus -symbol represents the bitwise exclusive-or operation (which is `^` in Python).

Polynomial Hash Codes

The summation and exclusive-or hash codes, described above, are not good choices for character strings or other sequences of characters. For example, consider a 16-bit hash code for a character string s that sums the Unicode values of the characters in s . This hash code unfortunately produces lots of unwanted collisions.

10.2. Hash Tables 411

10.2.1 Hash Functions

The goal of a hash function, h , is to map each key k to an integer in the range $[0, N-1]$, where N is the capacity of the bucket array for a hash table. Equipped with such a hash function, h , the main idea of this approach is to use the hash function value, $h(k)$, as an index into our bucket array, A , instead of the key k (which may not be appropriate for direct use as an index). That is, we store the item (k, v) in the bucket $A[h(k)]$.

If there are two or more keys with the same hash value, then two different items will be mapped to the same bucket in A . In this case, we say that a collision has occurred. To be sure, there are ways of dealing with collisions, which we will discuss later, but the best strategy is to be fast and easy to compute.

It is common to view the evaluation of a hash function, $h(k)$, as consisting of two portions—a hash code that maps a key k to an integer, and a compression function that maps the hash code to an integer within a range of indices, $[0, N-1]$, for a bucket array. (See Figure 10.5.)

-1 hash code

120 -2... ..

compression function

120N - 1 ...Arbitrary Objects

Figure 10.5: Two parts of a hash function: a hash code and a compression function.

The advantage of separating the hash function into two such components is that the hash code portion of that computation is independent of a specific hash table size. This allows the development of a general hash code for each object that can be used for a hash table of any size currently stored in the map. (See Section 10.2.3.)

Compression Functions

The hash code for a key k will typically not be suitable for immediate use with a bucket array, because the integer hash code may be negative or may exceed the capacity of the bucket array. Thus, once we have determined an integer hash code for a key object k , there is still the issue of mapping that integer into the range $[0, N-1]$. This computation, known as a compression function, is the second action performed as part of an overall hash function. A good compression function is one that minimizes the number of collisions for a given set of distinct hash codes.

The Division Method

A simple compression function is the division method, which maps an integer h to $h \bmod N$,

where N , the size of the bucket array, is a fixed positive integer. Additionally, if we take N to be a prime number, then this compression function helps spread out the distribution of hashed values. Indeed, if N is not prime, then there is greater risk

that patterns in the distribution of hash codes will be repeated in the distribution of hash values, thereby causing collisions. For example, if the hash codes $\{200, 205, 210, 215, 220, \dots, 600\}$ are mapped into a bucket array of size 100, then each hash code will collide with three others. But if we use a bucket array of size 101, then there will be no collisions. If a hash function is chosen well, it should ensure that the probability of two different keys having the same hash code is small. Choosing N to be a prime number is not always enough, however, for if there is a repeated pattern of hash codes of the form $pN + q$ for several different p 's, then there will still be collisions.

The MAD Method

A more sophisticated compression function, which helps eliminate repeated patterns in a set of integer keys, is the MAD method. This method maps an integer h to

$$[(a \cdot h + b) \bmod p] \bmod N,$$

where N is the size of the bucket array, p is a prime number larger than N , a and b are integers chosen at random from the interval $[0, p-1]$, with $a > 0$. This

compression function is chosen in order to eliminate repeated patterns in the set of hash codes and get us closer to the same as we would have if these keys were thrown into a uniformly at random.

10.2. Hash Tables 413

collisions for common groups of strings. In particular, "temp01" and "temp10" collide using this function, as do "stop", "tops", "pots", and "spot". A better hash code should somehow take into consideration the positions of the x_i 's. An alternative hash code, which does exactly this, is to choose a nonzero constant, a , and use as a hash code the value $x_0a^{n-1} + x_1a^{n-2} + \dots + x_{n-2}a + x_{n-1}$.

Mathematically speaking, this is simply a polynomial in a that takes the components $(x_0, x_1, \dots, x_{n-1})$ of an object x as its coefficients. This hash code is therefore called a polynomial hash code. By Horner's rule (see Exercise C-3.50), this polynomial can be computed as $x_{n-1} + a(x_{n-2} + a(x_{n-3} + \dots + a(x_2 + a(x_1 + ax_0)) \dots))$.

Intuitively, a polynomial hash code uses multiplication by different powers as a way to spread out the influence of each component across the resulting hash code.

Of course, on a typical computer, evaluating a polynomial will be done using the finite bit representation for a hash code; hence, the value will periodically overflow the bits used for an integer. One should be mindful that such overflows are occurring and choose the constant a so that it has some nonzero, low-order bits, which will serve to preserve some of the information content even as we overflow. We have done some experimental studies that suggest that 33, 37, 39, and 41 are particularly good choices for a when working with character strings that are English words. In fact, in a list of over 50,000 English words formed as the union of the word lists provided in two of the previous chapters, 33, 37, 39, or 41 produced less than 7 collisions in each case!

Cyclic-Shift Hash Codes

A variant of the polynomial hash code replaces multiplication by a with a cyclic shift of a partial sum by a certain number of bits. For example, a 5-bit cyclic shift of the 32-bit value 00111101100101101010100010101000 is achieved by taking the leftmost five bits and placing those on the rightmost side of the representation, resulting in 1011001011010101000101100101101000. While this operation has little natural meaning in terms of arithmetic, it accomplishes the goal of varying the bits of the calculation. In Python, a 32-bit integer is represented as a list of 32 bits.

10.2. Hash Tables 421

Operation

List

Hash Table

expected

worst case

getitem

$O(n)$

$O(1)$

$O(n)$

setitem

$O(n)$

$O(1)$

$O(n)$

delitem

$O(n)$

$O(1)$

$O(n)$

len

$O(1)$

$O(1)$

$O(1)$

iter

$O(n)$

$O(n)$

$O(n)$

Table 10.2: Comparison of the running times of the methods of a map realized by means of an unsorted list (as in Section 10.1.5) or a hash table. We let n denote the number of items in the map, and we assume that the bucket array supporting the hash table is maintained such that its capacity is proportional to the number of items in the map.

In practice, hash tables are among the most efficient means for implementing a map, and it is essentially taken for granted by programmers that their core operations run in constant time. Python's dict class is implemented with hashing, and the Python interpreter relies on dictionaries to retrieve an object that is referenced by an identifier in a given namespace. (See Sections 1.10 and 2.5.) The basic command `c=a+b` involves two calls to `getitem`

in the dictionary for the local

namespace to retrieve the values identified as `a` and `b`, and a call to

`setitem`

to store the result associated with name `c` in that namespace. In our own algorithm analysis, we simply presume that such dictionary operations run in constant time, independent of the number of entries in the dictionary.

In a 2003 academic paper [31], researchers discuss the possibility of exploiting a hash table's worst-case performance to cause a denial-of-service (DoS) attack of Internet technologies. For many years, they note that an attacker could precompute a very large number of moderate-length strings that all hash to the identical 32-bit hash code. (Recall that by any of the hashing schemes we describe, other than separate chaining, the hash code is used to index into the bucket array.)

In late 2011, another team of researchers demonstrated an implementation of just such an attack [61]. Web servers allow a series of key-value parameters to be embedded in a URL using a syntax like `http://www.example.com/?key=value&key2=value2`.

Typically, those key-value pairs are immediately stored in a map by the server, and a limit is placed on the length and number of such parameters presuming that storage time in the map will be linear in the number of entries. If all keys were to collide, that storage requires quadratic time in the number of entries.

10.2.3 Load Factors, Rehashing, and Efficiency

In the hash table schemes described thus far, it is important that the load factor, $\alpha = n/N$, be kept below 1. With separate chaining, as α gets very close to 1, the probability of a collision greatly increases, which adds overhead to our operations, since we must revert to linear-time list-based methods in buckets that have collisions.

Experiments and average-case analyses suggest that we should maintain $\alpha < 0.9$ for hash tables with separate chaining.

With open addressing, on the other hand, as the load factor α grows beyond 0.5

and starts approaching 1, clusters of entries in the bucket array start to grow as well. These clusters cause the probability of a collision to increase significantly.

Experiments suggest that we should maintain $\alpha < 0.5$ for an open addressing scheme with linear

probing, and perhaps only a bit higher for other open addressing schemes (for example, Python's implementation).

If an insertion causes the load factor of a hash table to go above the specified

threshold, then it is common to resize the table (to regain the specified load factor)

and to reinsert all objects into this new table. Although we need not define a new

hash code for each object, we do need to reapply a new compression function that takes into consideration the size of the new bucket array.

When rehashing to a new table, it is a good requirement for the new hash function to

scatter the items throughout the new bucket array. Indeed, if we always double the size of the table with each rehashing operation, then

we can amortize the cost of rehashing all the entries in the table against the time used to insert them in the first place.

Efficiency of Hash Tables

Although the details of the average-case analysis of hashing are beyond the scope of this book, its probabilistic behavior is well understood.

Thus, to store n entries, the expected number of keys in a bucket would

be n/N , which is $O(1)$ if n is $O(N)$.

The costs associated with a periodic rehashing, to resize a table after occasional

insertions or deletions can be accounted for separately, leading to an additional

$O(1)$ amortized cost for

setitem

and

getitem

.

In the worst case, a poor hash function could map every item to the same bucket.

This would result in linear-time performance for the core map operations with separate

chaining, or with any open addressing model in which the secondary sequence of probes depends only on the primary hash value.

Table 10.2.

452 Chapter 10. Maps, Hash Tables, and Skip Lists

10.6 Exercises

For help with exercises, please visit the site, www.wiley.com/college/goodrich.

Reinforcement

R-10.1 Give a concrete implementation of the `popmethod` in the context of the `MutableMapping` class, relying only on the `·ve` primary abstract methods of that class.

R-10.2 Give a concrete implementation of the `items()` method in the context of the `MutableMapping` class, relying only on the `·ve` primary abstract methods of that class. What would its running time be if directly applied to the `UnsortedTableMap` subclass?

R-10.3 Give a concrete implementation of the `items()` method directly within the `UnsortedTableMap` class, ensuring that the entire iteration runs in $O(n)$ time.

R-10.4 What is the worst-case running time for inserting n key-value pairs into an initially empty map M that is implemented with the `UnsortedTableMap` class?

R-10.5 Reimplement the `UnsortedTableMap` class from Section 10.1.5, using the `PositionalList` class from Section 7.4 rather than a Python list.

R-10.6 Which of the hash table collision-handling schemes could tolerate a load factor above 1 and which could not?

R-10.7 Our `Position` classes for lists and trees support the `eq` method so that

two distinct position instances are considered equivalent if they refer to the same underlying node in a structure.

hash

method that

is consistent with this notion of equivalence. Provide such a

hash

method.

R-10.8 What would be a good hash code for a vehicle identification number that is a string of numbers and letters?

R-10.9 Draw the 11-entry hash table that results from using the hash function, $h(i) = (3i + 5) \bmod 11$, to hash the keys 16, and 5, assuming collisions are handled by chaining.

R-10.10 What is the result of the previous exercise, assuming collisions are handled by linear probing?

R-10.11 Show the result of Exercise R-10.9, assuming collisions are handled by quadratic probing, up to the point where the method fails.

10.2 Hash Tables

In this section, we introduce one of the most practical data structures for implementing a map, and the one that is used by Python's own implementation of the dict class. This structure is known as a hash table.

Intuitively, a map M supports the abstraction of using keys as indices with a syntax such as $M[k]$. As a mental warm-up, consider a restricted setting in which a map with n items uses keys that are known to be integers in a range from 0 to $N-1$ for some N . In this case, we can represent the map using a lookup table of length N , as diagrammed in Figure 10.3.

```
0 123456789 10
D Z C Q
```

Figure 10.3: A lookup table with length 11 for a map containing items (1,D), (3,Z), (6,C), and (7,Q).

In this representation, we store the value associated with key k at index k of the table (presuming that we have a distinct way to represent an empty slot). Basic map operations of

```
,
getitem
, and
delitem
can be implemented in
O(1) worst-case time.
```

There are two challenges in extending this framework to the more general setting of a map. First, we may not wish to devote an array of length N if it is the case that $N \gg n$. Second, we do not in general require that a map's keys be integers.

The novel concept for a hash table is the use of a hash function to map general keys to corresponding indices in a table. Ideally, keys will be well distributed in the range from 0 to $N-1$ by a hash function that maps distinct keys to distinct indices. As a result, we will conceptualize our table as a bucket array, as shown in Figure 10.4, in which each bucket may manage a collection of items that are sent to a specific index by the hash function.

(To save space, an empty bucket may be replaced by None.)

```
0 123456789 10
(1,D) (25,C)
(3,F)
(14,Z) (39,C) (6,A) (7,Q)
```

Figure 10.4: A bucket array of capacity 11 with items (1,D), (25,C), (3,F), (14,Z), (6,A), (39,C), and (7,Q), using a simple hash function.

10.2. Hash Tables 417

10.2.2 Collision-Handling Schemes

The main idea of a hash table is to take a bucket array, A , and a hash function, h , and use them to implement a map by storing each item (k, v) in the bucket $A[h(k)]$.

This simple idea is challenged, however, when we have two distinct keys, k_1 and k_2 , such that $h(k_1) = h(k_2)$. The existence of such collisions prevents us from simply inserting a new item (k, v) directly into the bucket $A[h(k)]$. It also complicates our procedure for performing insertion, search, and deletion operations.

Separate Chaining

A simple and efficient way for dealing with collisions is to have each bucket $A[j]$ store its own secondary container, holding items (k, v) such that $h(k) = j$. A natural choice for the secondary container is a small map instance implemented using a list, as described in Section 10.1.5. This collision resolution rule is known as separate chaining, and is illustrated in Figure 10.6.

A 123456789 10 01 112

123825

9054

28413618 10

Figure 10.6: A hash table of size 13, storing 10 items with integer keys, with collisions resolved by separate chaining. The compression function is $h(k) = k \bmod 13$.

For simplicity, we do not show the values associated with the keys.

In the worst case, operations on an individual bucket take time proportional to the size of the bucket. Assuming we use a good hash function to index the items of our map in a bucket array of capacity N , the expected size of a bucket is n/N .

Therefore, if given a good hash function, the core map operations run in $O(n/N)$.

The ratio $\alpha = n/N$, called the load factor of the hash table, should be bounded by a small constant, preferably below 1. As long as α is $O(1)$, the core operations on the hash table run in $O(1)$ expected time.

10.2. Hash Tables 419

To implement a deletion, we cannot simply remove a found item from its slot in the array. For example, after the insertion of key 15 portrayed in Figure 10.7, if the item with key 37 were trivially deleted, a subsequent search for 15 would fail because that search would start at this special marker possibly occupying spaces in our hash table, we modify our search algorithm so that the search for a key will skip over cells containing the available marker and continue probing until reaching the desired item or an empty bucket (or returning back to where we started for

set item

should remember an available cell encountered during the search for a key, since this is a valid place to put a new item (k, v) , if no existing item is found.

Although use of an open addressing scheme can save space, linear probing suffers from an additional disadvantage. It tends to cluster the items of a map into contiguous runs, which may eventually fill the hash table (i.e., all cells in the hash table are occupied). Such contiguous runs of occupied hash cells cause searches to slow down considerably.

Another open addressing strategy, known as quadratic probing, iteratively tries the buckets $A[(h(k) + f(i)) \bmod N]$, for $i = 0, 1, 2, \dots$, where $f(i) = i^2$.

2, until finding

an empty bucket. As with linear probing, the quadratic probing strategy complicates the removal operation, but it does not cause clustering. When N is prime and the bucket strategy is guaranteed to find an empty slot. However, this guarantee is not valid once the table becomes at least half full. We explore the cause of this type of clustering in an exercise (C-10.36).

An open addressing strategy that does not cause clustering of the kind produced by linear probing or the kind produced by quadratic probing is the double hashing strategy. In this approach, we choose a secondary hash function, h_2 ,

where h_2 is a hash function

some key k to a bucket $A[h(k)]$ that is already occupied, then we iteratively try the buckets $A[(h(k) + f(i)) \bmod N]$ next, for $i = 1, 2, 3, \dots$, where $f(i) = i \cdot h_2(k)$.

In this scheme, the secondary hash function is not allowed to evaluate to zero; a common choice is $h_2(k) = q \cdot (k \bmod q)$, for some prime number $q < N$. Also, N should be a prime.

Another approach to avoid clustering with open addressing is to iteratively try buckets $A[(h(k) + f(i)) \bmod N]$ where $f(i)$ is based on a pseudo-random number generator, providing a repeatable, but somewhat arbitrary, sequence of subsequent probes that depends upon the seed value. This approach is currently used by Python's dictionary class.

10.2.4 Python Hash Table Implementation

In this section, we develop two implementations of a hash table, one using separate chaining and the other using open addressing with linear probing. While these approaches to collision resolution are different, they share a common class (from Code Fragment 10.2), to define a new HashMapBase class (see Code Fragment 10.4), providing much of the common functionality to our two hash table implementations. The main design goals are:

- The bucket array is represented as a Python list, named `self.buckets`.

- Each bucket is represented as a Python list, named `self.buckets[j]`, with all

- entries initialized to `None`.

- We maintain an instance variable `self.size`.

- `self.size` represents the number of distinct

- items that are currently stored in the hash table.

- If the load factor of the table increases beyond 0.5, we double the size of the table and rehash all items into the new table.

- We define a

- `hash`

- function utility method that relies on Python's built-in

- `hash` function to produce hash codes for keys, and a randomized Multiply-Add-and-Divide (MAD) formula for the compression function.

What is not implemented in the base class is any notion of how a bucket

should be represented. With separate chaining, each bucket will be an independent structure. With open addressing

In our design, the HashMapBase class presumes the following to be abstract methods, which must be implemented by each concrete subclass:

-

- bucket

- `getitem(j, k)`

- This method should search bucket `j` for an item having key `k`, returning the associated value, if found, or else raising a `KeyError`.

-

- bucket

- `setitem(j, k, v)`

- This method should modify bucket `j` so that key `k` becomes associated with value `v`. If the key already exists, the new value overwrites the existing value. Otherwise, a new item is inserted and this method is responsible for incrementing `self.size`.

-

-

- bucket

- `delitem(j, k)`

- This method should remove the item from bucket `j` having key `k`, or raise a `KeyError` if no such item exists. (`self.size` is decremented after this method.)

-

- `iter`

This is the standard map method to iterate through all keys of the map. Our base class does not delegate this on

414 Chapter 10. Maps, Hash Tables, and Skip Lists

An implementation of a cyclic-shift hash code computation for a character string in Python appears as follows:

```
def hash  
code(s):  
mask = (1 << 32) - 1 # limit to 32-bit integers  
h = 0  
for character in s:  
h = (h << 5 & mask) | (h >> 27) # 5-bit cyclic shift of running sum  
h += ord(character) # add in value of next character  
return h
```

As with the traditional polynomial hash code, fine-tuning is required when using a cyclic-shift hash code, as we must wisely choose the amount to shift by for each new character. Our choice of a 5 shift amounts (see Table 10.1).

Collisions

Shift

Total

Max

0	234735
1	623
165076	43
2	38471
13	3
3	7174
5	4
4	1379
3	5
5	190
3	6
6	502
2	7
7	560
2	8
8	5546
4	9
9	393
3	10
10	5194

10.6. Exercises 455

C-10.32 Perform experiments on our ChainHashMap and ProbeHashMap classes to measure its efficiency using random key sets and varying limits on the load factor (see Exercise R-10.15).

C-10.33 Our implementation of separate chaining in ChainHashMap conserves memory by representing empty buckets in the table as None, rather than as empty instances of a secondary structure. Because many of these buckets will hold a single item, a better option would be to have the table directly reference the Item instance, and to reserve use of secondary containers for buckets that have two or more items. Modify our implementation to provide this additional optimization.

C-10.34 Computing a hash code can be expensive, especially for lengthy keys. In our hash table implementation, inserting an item, and recompute each item's hash code each time we resize our table. Python's dict class makes an interesting trade-off. The hash code is computed once, when an item is inserted, and the hash code is stored as an extra field of the item component.

C-10.35 Describe how to perform a removal from a hash table that uses linear probing to resolve collisions where we do not use a special marker to represent deleted elements. That is, we must rearrange the contents so that it appears that the removed entry was never there.

C-10.36 The quadratic probing strategy has a clustering problem related to the way it looks for open slots. Namely, it checks buckets $A[(h(k) + i^2) \bmod N]$, for $i = 1, 2, \dots, N-1$.

a. Show that $i^2 \bmod N$ will assume at most $(N+1)/2$ distinct values, for N prime, as i ranges from 1 to $N-1$. As a part of this justification, note that $i^2 \bmod N = (N-i)^2 \bmod N$ for all i .

b. A better strategy is to choose a prime N such that $N \bmod 4 = 3$ and then to check the buckets $A[(h(k) \pm i^2) \bmod N]$ as i ranges from 1 to $(N-1)/2$, alternating between plus and minus. Show that this alternate version is guaranteed to check every bucket in A .

C-10.37 Refactor our ProbeHashMap design so that the sequence of secondary probes for collision resolution can be more easily customized. Demonstrate your new framework by providing several examples.

C-10.38 Design a variation of binary search for performing the multimap operation $\text{nd_all}(k)$ implemented with a sorted search table that includes duplicates, and show that it runs in time $O(s + \log n)$, where n is the number of elements in the dictionary and s is the number of items with given key k .

10.2. Hash Tables 423

1class HashMapBase(MapBase):

2...Abstract base class for map using hash-table with MAD compression....

3

4def

init

(self, cap=11, p=109345121):

5 ...Create an empty hash-table map....

6 self.

table = cap

[None]

7 self.

n=0 # number of entries in the map

8 self.

prime = p # prime for MAD compression

9 self.

scale = 1 + randrange(p - 1) # scale factor for MAD

10 self.

shift = randrange(p) # shift from 0 to p-1 for MAD

11

12def

hash

function(self,k) :

13 return (hash(k)

self.

scale + self.

shift) % self.

prime % len(self.

table)

14

15def

len

(self):

16 return self .

n

1718def

getitem

(self,k) :

19 j=self.

hash

function(k)

20 return self .

bucket

getitem(j, k) # may raise KeyError

21

22def

setitem

(self,k ,v) :

23 j=self.

hash

10.6. Exercises 453

R-10.12 What is the result of Exercise R-10.9 when collisions are handled by double hashing using the secondary hash function $h_{\text{prime}}(k) = 7 \cdot (k \bmod 7)$?

R-10.13 What is the worst-case time for putting n entries in an initially empty hash table, with collisions resolved by chaining? What is the best case?

R-10.14 Show the result of rehashing the hash table shown in Figure 10.6 into a table of size 19 using the new hash function $h(k) = 3k \bmod 17$.

R-10.15 Our `HashMapBase` class maintains a load factor ≈ 0.5 . Reimplement that class to allow the user to specify the maximum load, and adjust the concrete subclasses accordingly.

R-10.16 Give a pseudo-code description of an insertion into a hash table that uses quadratic probing to resolve collisions, assuming we also use the trick of replacing deleted entries with a special marker.

R-10.17 Modify our `ProbeHashMap` to use quadratic probing.

R-10.18 Explain why a hash table is not suited to implement a sorted map.

R-10.19 Describe how a sorted list implemented as a doubly linked list could be used to implement the sorted map.

R-10.20 What is the worst-case asymptotic running time for performing n deletions from a `SortedTableMap` instance that initially contains $2n$ entries?

R-10.21 Consider the following variant of the

```
.nd
index method from Code Fragment 10.8, in the context of the SortedTableMap class:
```

```
def
.nd
index(self, k, low, high):
    if high < low:
        return high + 1
    else:
        mid = (low + high) // 2
        if self.table[mid].key < k:
            return self.index(k, mid + 1, high)
        else:
            return self.index(k, low, mid - 1)
```

Does this always produce the same result as the original version? Justify your answer.

R-10.22 What is the expected running time of the methods for maintaining a max-ima set if we insert n pairs such that each pair has lower cost and performance than one before it? What is contained in the sorted map at the end of this series of operations? What if each pair had a lower cost and higher performance than the one before it?

R-10.23 Draw an example skip list S that results from performing the following series of operations on the skip list shown in Figure 10.13: `del S[38]`, `S[48] =`

```
x
, S[24] =
y
```


10.6. Exercises 457

C-10.49 Python's collections module provides an `OrderedDict` class that is unrelated to our sorted map abstraction. An `OrderedDict` is a subclass of the standard `hash`-based dictclass that retains the expected $O(1)$ performance for the primary map operations, but that also guarantees that the `iter`

method reports items of the map according to `·rst-in`, `·rst-out` (FIFO) order. That is, the key that has been in the dictionary the longest is reported `·rst`. (The order is unaffected when the value for an existing key is overwritten.) Describe an algorithmic approach for achieving such performance.

Projects

P-10.50 Perform a comparative analysis that studies the collision rates for various

`hash` codes for character strings, such as various polynomial `hash` codes

for different values of the parameter `a`. Use a `hash` table to determine

collisions, but only count collisions where different strings map to the same `hash` code (not if they map to the same

Test these `hash` codes on text files found on the Internet.

P-10.51 Perform a comparative analysis as in the previous exercise, but for 10-digit telephone numbers instead of character strings.

P-10.52 Implement an `OrderedDict` class, as described in Exercise C-10.49, ensuring that the primary map operations run in $O(1)$ expected time.

P-10.53 Design a Python class that implements the skip-list data structure. Use this class to create a complete implementation of the sorted map ADT.

P-10.54 Extend the previous project by providing a graphical animation of the skip-list operations. Visualize how entries move up the skip list during insertions and are linked out of the skip list

P-10.55 Write a spell-checker class that stores a lexicon of words, `W`, in a Python

set, and implements a method, `check(s)`, which performs a spell check

on the string `s` with respect to the set of words, `W`. If `s` is in `W`, then

the call to `check(s)` returns a list containing only `s`, as it is assumed to

be spelled correctly in this case. If `s` is not in `W`, then the call to `check(s)`

returns a list of every word in `W` that might be a correct spelling of `s`. Your

program should be able to handle all the common ways that `s` might be a

misspelling of a word in `W`, including swapping adjacent characters in a

word, inserting a single character in between two adjacent characters in a word, deleting a single character from

a word with another character. For an extra challenge, consider phonetic

substitutions as well.

418 Chapter 10. Maps, Hash Tables, and Skip Lists

Open Addressing

The separate chaining rule has many nice properties, such as affording simple implementations of map operations, but it nevertheless has one slight disadvantage:

It requires the use of an auxiliary data structure—a list—to hold items with colliding keys. If space is at a premium (as it often is), these structures are employed, but it requires a bit more complexity to deal with collisions. There are several variants of this approach, collectively referred to as open addressing schemes, which we discuss next. Open addressing requires that the load factor is always at most 1 and that items are stored directly in the cells of the bucket array itself.

Linear Probing and Its Variants

A simple method for collision handling with open addressing is linear probing.

With this approach, if we try to insert an item (k, v) into a bucket $A[j]$ that is already occupied, where $j = h(k)$, then we next try $A[(j+1) \bmod N]$. If $A[(j+1) \bmod N]$

is also occupied, then we try $A[(j+2) \bmod N]$, and so on, until we find an empty

bucket that can accept the new item. Once this bucket is located, we simply in-

sert the item there. Of course, this collision resolution strategy requires that we change the implementation when

getitem

, or

delitem

operations. In particular, to attempt

to locate an item with key equal to k , we must examine consecutive slots, starting

from $A[h(k)]$, until we either find an item with that key or we find an empty bucket.

(See Figure 10.7.) The name “linear probing” comes from the fact that accessing a cell of the bucket array can be

26123456789 1 0 0 New element with

key = 15 to be inserted Must probe 4 times

before finding empty slot

53 7 1 6 2 1 13

Figure 10.7: Insertion into a hash table with integer keys using linear probing. The

hash function is $h(k) = k \bmod 11$. Values associated with keys are not shown.

424 Chapter 10. Maps, Hash Tables, and Skip Lists

Separate Chaining

Code Fragment 10.5 provides a concrete implementation of a hash table with separate chaining, in the form of the ChainHashMap class. To represent a single bucket, it relies on an instance of the UnsortedTableMap class from Code Fragment 10.3.

The first three methods in the class use index j to access the potential bucket in the bucket array, and a check for the special case in which that table entry is None.

The only time we need a new bucket structure is when

bucket

setitem is called on

an otherwise empty slot. The remaining functionality relies on map behaviors that are already supported by the inbuilt dict class. A small bit of forethought to determine whether the application of

setitem

on the chain

causes a net increase in the size of the map (that is, whether the given key is new).

1 class ChainHashMap(HashMapBase):

2 """Hash map implemented with separate chaining for collision resolution. ...

34 def

bucket

getitem(self, j, k) :

5 bucket = self.

table[j]

6 if bucket is None :

7 raise KeyError(

Key Error:

+repr(k)) # no match found

8 return bucket[k] # may raise KeyError

9

10 def

bucket

setitem(self, j, k, v) :

11 if self.

table[j] is None :

12 self.

table[j] = UnsortedTableMap() # bucket is new to the table

13 oldsize = len(self.

table[j])

14 self.

table[j][k] = v

15 if len(self.

table[j]) > oldsize: # key was new to the table

16 self.

n += 1 # increase overall map size

17 18 def

bucket

delitem(self, j, k) :

19 bucket = self.

table[j]

20 if bucket is None :

21 raise KeyError(

742 Index

directed, 620, 621, 657

acyclic, 655-657

strongly connected, 623

mixed, 621

reachability, 651-654

shortest paths, 654simple, 622traversal, 638-650

undirected, 620, 621

weighted, 659-696

greedy method, 603, 660, 661Guibas, Leonidas, 535

Gutttag, John, 108, 254, 298

Harmonic number, 131, 180, 728

hash code, 411-415

cyclic-shift, 413-414

polynomial, 413

hash function, 415

hash table, 410-426

clustering, 419

collision, 411

collision resolution, 417-419

double hashing, 419

linear probing, 418

quadratic probing, 419

HashMapBase class, 422-423

header sentinel, 270heap, 370-384

bottom-up construction, 380-384

heap-sort, 384, 388-389

HeapPriorityQueue class, 377-378, 382

heapq module, 384

height of a tree, 309-310, 474

height-balance property, 481, 483

Hell, Pavol, 696Hennessy, John, 719heuristic, 289hierarchy, 82

Hoare, C. A. R., 580

hook, 342, 468, 478Hopcroft, John, 254, 298, 535, 696Hopper, Grace, 36

Horner-s method, 146

HTML, 236-238, 251, 582Huang, Bing-Chao, 580Huffman coding, 601-602I/O complexity, 711idfunction, 29

identifier, 41DLE, 2

immutable type, 7, 11, 415

implied method, 76

import statement, 48

inoperator, 14-15, 75

in-degree, 621in-place algorithm, 389, 396, 559, 702

incidence collection, 630incident, 621

incoming edges, 621

independent, 729, 730

index, 186

negative, 14

zero-indexing, 9, 14

IndexError, 20, 33, 34, 83, 232, 303

induction, 138-139, 162

Contents xv

8.4.1 Preorder and Postorder Traversals of General Trees . . .	328
8.4.2 Breadth-First Tree Traversal	330
8.4.3 Inorder Traversal of a Binary Tree	331
8.4.4 Implementing Tree Traversals in Python	333
8.4.5 Applications of Tree Traversals	337
8.4.6 Euler Tours and the Template Method Pattern	341
8.5 Case Study: An Expression Tree	348
8.6 Exercises	352
9 Priority Queues	362
9.1 The Priority Queue Abstract Data Type	363
9.1.1 Priorities	363
9.1.2 The Priority Queue ADT	364
9.2 Implementing a Priority Queue	365
9.2.1 The Composition Design Pattern	365
9.2.2 Implementation with an Unsorted List	366
9.2.3 Implementation with a Sorted List	368
9.3 Heaps	370
9.3.1 The Heap Data Structure	370
9.3.2 Implementing a Priority Queue with a Heap	372
9.3.3 Array-Based Representation of a Complete Binary Tree .	376
9.3.4 Python Heap Implementation	376
9.3.5 Analysis of a Heap-Based Priority Queue	379
9.3.6 Bottom-Up Heap Construction	380
9.3.7 Python's <code>heapq</code> Module	384
9.4 Sorting with a Priority Queue	385
9.4.1 Selection-Sort and Insertion-Sort	386
9.4.2 Heap-Sort	388
9.5 Adaptable Priority Queues	390
9.5.1 Locators	390
9.5.2 Implementing an Adaptable Priority Queue	391
9.6 Exercises	395
10 Maps, Hash Tables, and Skip Lists	401
10.1 Maps and Dictionaries	402
10.1.1 The Map ADT	403
10.1.2 Application: Counting Word Frequencies	405
10.1.3 Python's <code>MutableMapping</code> Abstract Base Class	406
10.1.4 Our <code>Map</code> Base Class	407
10.1.5 Simple Unsorted Map Implementation	408
10.2 Hash Tables	410
10.2.1 Hash Functions	411

Chapter

10 Maps, Hash Tables, and Skip Lists

Contents

10.1 Maps and Dictionaries	402
10.1.1 The Map ADT	403
10.1.2 Application: Counting Word Frequencies	405
10.1.3 Python's MutableMapping Abstract Base Class	406
10.1.4 Our Map Base Class	407
10.1.5 Simple Unsorted Map Implementation	408
10.2 Hash Tables	410
10.2.1 Hash Functions	411
10.2.2 Collision-Handling Schemes	412
10.2.3 Load Factors, Rehashing, and Efficiency	420
10.2.4 Python Hash Table Implementation	422
10.3 Sorted Maps	427
10.3.1 Sorted Search Tables	428
10.3.2 Two Applications of Sorted Maps	434
10.4 Skip Lists	437
10.4.1 Search and Update Operations in a Skip List	439
10.4.2 Probabilistic Analysis of Skip Lists	443
10.5 Sets, Multisets, and Multimaps	446
10.5.1 The Set ADT	446
10.5.2 Python's MutableSet Abstract Base Class	448
10.5.3 Implementing Sets, Multisets, and Multimaps	450
10.6 Exercises	452

10.2. Hash Tables 425

Linear Probing

Our implementation of a `ProbeHashMap` class, using open addressing with linear probing, is given in Code Fragments 10.6 and 10.7. In order to support deletions, we use a technique described in Section 10.2.2 in which we place a special marker in a table location at which an item has been deleted, so that we can distinguish between it and a location that has `AVAIL`, as a sentinel. (We use an instance of the built-in object class because we do not care about any behaviors of the sentinel, just our ability to differentiate it from other objects.)

The most challenging aspect of open addressing is to properly trace the series of probes when collisions occur during an insertion or search for an item. To this end, we define a nonpublic utility method

`slot`, that searches for an item with key `k`

in `bucket[j]` (that is, where `j` is the index returned by the `hash` function for key `k`).

```
1 class ProbeHashMap(HashMapBase):
```

```
2     """Hash map implemented with linear probing for collision resolution...."""
```

```
3
```

```
    AVAIL = object() # sentinel marks locations of previous deletions
```

```
45 def
```

```
    is
```

```
    available( self, j ) :
```

```
6     """Return True if index j is available in table...."""
```

```
7     return self .
```

```
    table[j] is None or self .
```

```
    table[j] is not ProbeHashMap.
```

```
    AVAIL
```

```
89 def
```

```
    .nd
```

```
    slot(self, j, k) :
```

```
10    """Search for key k in bucket at index j.
```

```
11 12 Return (success, index) tuple, described as follows:
```

```
13 If match was found, success is True and index denotes its location.
```

```
14 If no match found, success is False and index denotes first available slot.
```

```
15 ...
```

```
16 .rstAvail = None
```

```
17 while True :
```

```
18 if self.
```

```
    is
```

```
    available(j):
```

```
19 if .rstAvail is None :
```

```
20 .rstAvail = j # mark this as first avail
```

```
21 if self.
```

```
    table[j] is None :
```

```
22 return (False, .rstAvail) # search has failed
```

```
23 elif k == self.
```

```
    table[j].
```

```
    key:
```

```
24 return (True, j) # found a match
```

```
25 j = ( j + 1 ) % len ( self.
```

```
    table) # keep looking (cyclically)
```

10.3.1 Sorted Search Tables

Several data structures can efficiently support the sorted map ADT, and we will examine some advanced techniques in Section 10.4 and Chapter 11. In this section, we begin by exploring a simple realization of a sorted map, where the keys are stored in an array in sorted order. We refer to this implementation of a map as a sorted search table.

```
9 2 4 5 7 8 12 14 17 19 22 25 27 28 33 5
3 7 0 1 2 3 4 6 7 8 9 1 0 1 1 1 2 1 3 1 4 1 5
```

Figure 10.8: Realization of a map by means of a sorted search table. We show only the keys for this map, so as to highlight their ordering.

As was the case with the unsorted table map of Section 10.1.5, the sorted search table has a space requirement that is $O(n)$, assuming we grow and shrink the array to keep its size proportional to the number of items in the map. The primary advantage of this representation, and our reason for insisting that it be array-based, is that it allows us to use the binary search algorithm for a variety of efficient operations.

Binary Search and Inexact Searches

We originally presented the binary search algorithm in Section 4.1.3, as a means for detecting whether a given target is stored within a sorted sequence. In our original presentation (Code Fragment 4.3 on page 156), a binary search function returned

True or False to designate whether the desired target was found. While such an approach could be used to implement the contains method of the map ADT,

we can adapt the binary search algorithm to provide far more useful information when performing forms of inexact search in support of the sorted map ADT.

The important realization is that while performing a binary search, we can determine the index at or near where a target might be found. During a successful search, the standard implementation determines the precise index at which the target is found. During an unsuccessful search, although the target is not found, the algorithm will effectively determine a pair of indices designating elements of the collection that are just less than and just greater than the target.

As a motivating example, our original simulation from Figure 4.5 on page 156 shows a successful binary search for a target of 22, using the same data we portray in Figure 10.8. Had we instead been searching for a target of 20, we would have determined that the missing target lies in the gap between values 19 and 22 in that example.

446 Chapter 10. Maps, Hash Tables, and Skip Lists

10.5 Sets, Multisets, and Multimaps

We conclude this chapter by examining several additional abstractions that are closely related to the map ADT, and that can be implemented using data structures similar to those for a map.

- A set is an unordered collection of elements, without duplicates, that typically supports efficient membership tests. In essence, elements of a set are like keys of a map, but without any associated values.
- A multiset (also known as a bag) is a set-like container that allows duplicates.
- A multimap is similar to a traditional map, in that it associates values with keys; however, in a multimap the same key can be mapped to multiple values. For example, the index of this book maps a given term to one or more locations at which the term occurs elsewhere in the book.

10.5.1 The Set ADT

Python provides support for representing the mathematical notion of a set through the built-in classes `frozenset` and `set`, `frozenset` being an immutable form. Both of those classes are implemented using hash tables in Python.

Python's `collections` module defines abstract base classes that essentially mirror these built-in classes. Although the choice of names is counterintuitive, the abstract base class `collections.Set` mirrors the built-in `set` class, and the abstract base class `collections.MutableSet` is akin to the concrete `set` class.

In our own discussion, we equate the "set ADT" with the behavior of the built-in `set` class (and thus, the `collections.MutableSet` base class). We begin by listing what we consider to be the "most fundamental behaviors for a set `S`":

`S.add(e)` : Add element `e` to the set. This has no effect if the set already contains `e`.

`S.discard(e)` : Remove element `e` from the set, if present. This has no effect if the set does not contain `e`.

`e in S` : Return `True` if the set contains element `e`. In Python, this is implemented with the special `__contains__` method.

`len(S)` : Return the number of elements in set `S`. In Python, this is implemented with the special method `__len__`.

`iter(S)` : Generate an iteration of all elements of the set. In Python, this is implemented with the special method `__iter__`.

454 Chapter 10. Maps, Hash Tables, and Skip Lists

R-10.24 Give a pseudo-code description of the `delitem`

`map` operation when using a skip list.

R-10.25 Give a concrete implementation of the `pop` method, in the context of a `MutableSet` abstract base class, that relies only on the core set behaviors described in Section 10.5.2.

R-10.26 Give a concrete implementation of the `isdisjoint` method in the context of the `MutableSet` abstract base class, relying only on the primary abstract methods of that class. Your algorithm should run in $O(\min(n, m))$ where n and m denote the respective cardinalities of the two sets.

R-10.27 What abstraction would you use to manage a database of friends' birthdays in order to support efficient queries such as "find all friends whose birthday is today" and "find the friend who was born on this day"? Creativity

C-10.28 On page 406 of Section 10.1.3, we give an implementation of the method `setdefault` as it might appear in `UnsortedTableMap`. While that method accomplishes the goal in a general fashion, its efficiency is less than ideal. In particular, when `getitem`

`getitem` is called, and then a subsequent insertion via `setitem`

is performed, `setitem` must scan the entire table. For a concrete implementation, such as `UnsortedTableMap`, this is twice the work because a complete scan of the table will take place during the failed `getitem`, and then another complete scan of the table takes place due to the implementation of `setitem`.

A better solution is for the `UnsortedTableMap` class to override `setdefault` to provide a direct solution that performs a single search. Give such an implementation of `UnsortedTableMap.setdefault`.

C-10.29 Repeat Exercise C-10.28 for the `ProbeHashMap` class.

C-10.30 Repeat Exercise C-10.28 for the `ChainHashMap` class.

C-10.31 For an ideal compression function, the capacity of the bucket array for a hash table should be a prime number. Finding such a prime by using the sieve algorithm. In this algorithm, we allocate a 2M-cell Boolean array A , such that cell i is associated with the integer i . We then initialize the array cells to all be `True` and we mark off all the cells that are multiples of 2, 3, 5, 7, and so on. This process can stop after it reaches a number larger than $2M$. (Hint: Consider a bootstrapping method for finding the primes up to $2M$.)

A trie T for a set S of strings can be used to implement a set or map whose keys are the strings of S . Namely, we perform a search in T for a string X by tracing down from the root the path indicated by the characters in X . If this path can be traced and terminates at a leaf node, then we know X is a key in the map. For example, in the trie in Figure 13.10, tracing the path for `·bull·` ends up at a leaf. If the path cannot be traced or the path can be traced but terminates at an internal node, then X is not a key in the map. In the example in Figure 13.10, the path for `·bet·` cannot be traced and the path for `·be·` ends at an internal node. Neither such word is in the map.

It is easy to see that the running time of the search for a string of length m is $O(m \cdot |\Sigma|)$, because we visit at most $m+1$ nodes of T and we spend $O(|\Sigma|)$ time at each node determining the child having the subsequent character as a label. The $O(|\Sigma|)$ upper bound on the time to locate a child with a given label is achievable, even if the children of a node are unordered, since there are at most $|\Sigma|$ children. We can improve the time spent at a node to be $O(\log |\Sigma|)$ or expected $O(1)$, by mapping characters to children using a secondary search table or **hash** table at each node, or by using a direct lookup table of size $|\Sigma|$ at each node, if $|\Sigma|$ is sufficiently small (as is the case for DNA strings). For these reasons, we typically expect a search for a string of length m to run in $O(m)$ time.

From the discussion above, it follows that we can use a trie to perform a special type of pattern matching, called word matching, where we want to determine whether a given pattern matches one of the words of the text exactly. Word matching differs from standard pattern matching because the pattern cannot match an arbitrary substring of the text—only a word of the original document must be added to the trie. (See Figure 13.11.) A simple extension of this scheme supports pre- x -matching queries. However, arbitrary occurrences of the pattern are not supported. To construct a standard trie for a set S of strings, we can use an incremental algorithm that inserts the strings one at a time. Recall the assumption that no string of S is a pre- x of another string. To insert a string X into the current trie T , we trace the path associated with X in T , creating a new chain of nodes to store the remaining characters of X when we get stuck. The running time to insert X with length m is similar to a search, with worst-case $O(m \cdot |\Sigma|)$ performance, or expected $O(m)$ if using secondary **hash** tables at each node. Thus, constructing the entire trie for set S takes expected $O(n)$ time, where n is the total length of the strings of S . There is a potential space inefficiency in the standard trie that has prompted the development of the compressed trie, which is also known (for historical reasons) as the Patricia trie. Namely, there are potentially a lot of nodes in the standard trie that have only one child, and the existence of such nodes is a waste. We discuss the compressed trie next.

14.2. Data Structures for Graphs 635

```
1#----- nested Vertex class -----
2classVertex:
3    ...Lightweight vertex structure for a graph...
4
5    slots
6    =
7    _element
8
9    def
10    init
11    (self,x ) :
12    ...Do not call constructor directly. Use Graph
13    si n s e r t
14    vertex(x)...
15
16    self.
17    element = x
18
19
20    defelement( self):
21    ...Return element associated with this vertex...
22    return self .
23    element
24
25
26    def
27    hash
28    (self): # will allow vertex to be a map/set key
29
30    15 return hash(id( self))
31
321617 #----- nested Edge class -----
3318classEdge:
3419    ...Lightweight edge structure for a graph...
3520
36    slots
37    =
38    _origin
39    ,
40    _destination
41    ,
42    _element
43
442122 def
45    init
46    (self,u ,v ,x ) :
47    ...Do not call constructor directly. Use Graph
48    si n s e r t
49    edge(u,v,x)...
50
5124 self.
52    origin = u
53
5425 self.
55    destination = v
56
5726 self.
58    element = x
```

740 Index

Cormen, Thomas, 535, 696

Counter class, 450

CPU, 111

CRC cards, 63

CreditCard class, 63, 69-73, 73, 83-86

Crochemore, Maxime, 618

cryptography, 216-218ctypes module, 191, 195

cubic function, 119

cyber-dollar, 197-199, 497-500, 682

cycle, 623

directed, 623

cyclic-shift **hash** code, 413-414

DAG, seedirected acyclic graph

data packets, 227

data structure, 110

Dawson, Michael, 55

debugging, 62

decision tree, 311, 463, 562

decorate-sort-undecorate design pattern,
570

decrease-and-conquer, 571

decryption, 216

deep copy, 102, 188

deepcopy function, 102, 188

defkeyword, 23

degree of a vertex, 621

deloperator, 15, 75

DeMorgan-s Law, 137

Demurjian, Steven, 108, 254depth of a tree, 308-310depth-first search (DFS), 639-647

deque, 247-249

abstract data type, 247-248

linked-list implementation, 249, 275

deque class, 249, 251

descendant, 302design patterns, v, 61

adapter, 231

amortization, 197-200

brute force, 584

composition, 287, 365, 407divide-and-conquer, 538-542, 550-
551

dynamic programming, 594-600

factory method, 479greedy method, 603

position, 279-281

prune-and-search, 571-573

template method, 93, 342, 406, 448,
478

DFS, seedepth-first search

Di Battista, Giuseppe, 361, 696

diameter, 358dictclass, 7, 11, 402

dictionary, 11, 16, 402-408, see also map

Preface vii

Contents and Organization

The chapters for this book are organized to provide a pedagogical path that starts with the basics of Python programming and object-oriented design. We then add foundational techniques like algorithm analysis and recursion. In the main portion of the book, we present fundam

1. Python Primer

2. Object-Oriented Programming

3. Algorithm Analysis

4. Recursion

5. Array-Based Sequences

6. Stacks, Queues, and Deques

7. Linked Lists

8. Trees

9. Priority Queues

10. Maps, Hash Tables, and Skip Lists

11. Search Trees

12. Sorting and Selection

13. Text Processing

14. Graph Algorithms

15. Memory Management and B-Trees

A. Character Strings in Python

B. Useful Mathematical Facts

A more detailed table of contents follows this preface, beginning on page xi.

Prerequisites

We assume that the reader is at least vaguely familiar with a high-level programming language, such as C, C++

- Variables and expressions.

- Decision structures (such as if-statements and switch-statements).

- Iteration structures (for loops and while loops).

- Functions (whether stand-alone or object-oriented methods).

For readers who are familiar with these concepts, but not with how they are expressed in Python, we provide a p
give a comprehensive treatment of Python.

1.2. Objects in Python 11

The set and frozenset Classes

Python's set class represents the mathematical notion of a set, namely a collection of elements, without duplicates, and without an inherent order to those elements.

The major advantage of using a set, as opposed to a list, is that it has a highly

optimized method for checking whether a specific element is contained in the set. This is based on a data structure

(topic of Chapter 10). However, there are two important restrictions due to the

algorithmic underpinnings. The first is that the set does not maintain the elements

in any particular order. The second is that only instances of immutable types can be

added to a Python set. Therefore, objects such as integers, floating-point numbers,

and character strings are eligible to be elements of a set. It is possible to maintain a

set of tuples, but not a set of lists or a set of sets, as lists and sets are mutable. The

frozenset class is an immutable form of the set type, so it is legal to have a set of

frozensets.

Python uses curly braces {} as delimiters for a set, for example, as {1, 7}

or {

red

,

green

,

blue

}. The exception to this rule is that {} does not

represent an empty set; for historical reasons, it represents an empty dictionary (see next paragraph). Instead, the

If an iterable parameter is sent to the constructor, then the set of distinct elements

is produced. For example, set(

hello

) produces {

h

,

e

,

l

,

o

}.

The dict Class

Python's dict class represents a dictionary, a mapping, from a set of distinct keys

to associated values. For example, a dictionary might map from unique student ID

numbers, to larger student records (such as the student's name, address, and course

grades). Python implements a dict using an almost identical approach to that of a

set, but with storage of the associated values.

A dictionary literal also uses curly braces, and because dictionaries were intro-

duced in Python prior to sets, the literal form {} produces an empty dictionary.

A nonempty dictionary is expressed using a comma-separated series of key:value pairs. For example, the dictionary

ga

:

Irish

,

de

:

1.5. Functions 29

Common Built-In Functions

Calling Syntax

Description

`abs(x)`

Return the absolute value of a number.

`all(iterable)`

Return True if `bool(e)` is True for each element `e`.

`any(iterable)`

Return True if `bool(e)` is True for at least one element `e`.

`chr(integer)`

Return a one-character string with the given Unicode code point.

`divmod(x, y)`

Return `(x // y, x % y)` as tuple, if `x` and `y` are integers.

`hash(obj)`

Return an integer `hash` value for the object (see Chapter 10).

`id(obj)`

Return the unique integer serving as an `identity` for the object.

`input(prompt)`

Return a string from standard input; the prompt is optional.

`isinstance(obj, cls)`

Determine if `obj` is an instance of the class (or a subclass).

`iter(iterable)`

Return a new iterator object for the parameter (see Section 1.8).

`len(iterable)`

Return the number of elements in the given iteration.

`map(f, iter1, iter2, ...)`

Return an iterator yielding the result of function calls `f(e1, e2, ...)` for respective elements `e1`, `e2`, ...

`max(iterable)`

Return the largest element of the given iteration.

`max(a, b, c, ...)`

Return the largest of the arguments.

`min(iterable)`

Return the smallest element of the given iteration.

`min(a, b, c, ...)`

Return the smallest of the arguments.

`next(iterator)`

Return the next element reported by the iterator (see Section 1.8).

`open(filename, mode)`

Open a file with the given name and access mode.

`ord(char)`

Return the Unicode code point of the given character.

`pow(x, y)`

Return the value `xy` (as an integer if `x` and `y` are integers); equivalent to `x`

`y`.

`pow(x, y, z)`

Return the value `(x mod z)` as an integer.

`print(obj1, obj2, ...)`

50 Chapter 1. Python Primer

next number in a sequence based upon one or more past numbers that it has generated. Indeed, a simple yet popular pseudo-random number generator chooses its next number based solely on the most recently chosen number and some additional parameters using the following

$$\text{next} = (\text{a} * \text{current} + \text{b}) \% \text{n};$$

where a , b , a n d n are appropriately chosen integers. Python uses a more advanced technique known as a Mersenne twister. It turns out that the sequences generated by these techniques can be proven to be statistically uniform, which is usually good enough for most applications requiring random numbers, such as games. For applications, such as computer security settings, where one needs unpredictable random sequences, this kind of randomness comes from outer space.

Since the next number in a pseudo-random generator is determined by the previous number(s), such a generator always needs a place to start, which is called its seed. The sequence of numbers generated for a given seed will always be the same.

One common trick to get a different sequence each time a program is run is to use a seed that will be different for each run. For example, we could use some timed input from a user or the current system time in milliseconds.

Python's random module provides support for pseudo-random number generation by defining a Random class; instances of that class serve as generators with independent state. This allows different aspects of a program to rely on their own pseudo-random number generators. The random module (essentially using a single generator instance for all top-level calls).

Syntax

Description

`seed(hashable)`

Initializes the pseudo-random number generator based upon the `hash` value of the parameter `random()`

Returns a pseudo-random floating-point value in the interval `[0.0,1.0)`.

`randint(a,b)`

Returns a pseudo-random integer in the closed interval `[a,b]`.

`randrange(start, stop, step)`

Returns a pseudo-random integer in the standard Python range indicated by the parameters.

`choice(seq)`

Returns an element of the given sequence chosen pseudo-randomly.

`shuffle(seq)`

Reorders the elements of the given sequence pseudo-randomly.

Table 1.8: Methods supported by instances of the Random class, and as top-level functions of the random module.

2.3. Class Definitions 75

Common Syntax

Special Method Form

$a+b$

$a.$

`add`

(b) ; alternatively $b.$

`radd`

(a)

$a \cdot b$

$a.$

`sub`

(b) ; alternatively $b.$

`rsub`

(a)

a

b

$a.$

`mul`

(b) ; alternatively $b.$

`rmul`

(a)

a/b

$a.$

`truediv`

(b) ; alternatively $b.$

`rtruediv`

(a)

$a//b$

$a.$

`floordiv`

(b) ; alternatively $b.$

`rfloordiv`

(a)

$a \% b$

$a.$

`mod`

(b) ; alternatively $b.$

`rmod`

(a)

a

b

$a.$

`pow`

(b) ; alternatively $b.$

`rpow`

(a)

$a \ll b$

$a.$

`lshift`

10.1 Maps and Dictionaries

Python's dict class is arguably the most significant data structure in the language. It represents an abstraction known as a dictionary in which unique keys are mapped to associated values. Because of the relationship they express between keys and values, dictionaries are commonly known as associative arrays or maps. In this book, we use the term dictionary when specifically discussing Python's dict class, and the term map when discussing the more general notion of the abstract data type. As a simple example, Figure 10.1 illustrates a map from the names of countries to their associated units of currency.

Rupee	Turkey	Spain	China	United States	India	Greece
Lira	Euro	Yuan	Dollar			

Figure 10.1: A map from countries (the keys) to their units of currency (the values).

We note that the keys (the country names) are assumed to be unique, but the values (the currency units) are not necessarily unique. For example, we note that Spain and Greece both use the euro for currency. Maps use an array-like syntax for indexing, such as `currency[Greece]` to access a value associated with a given key.

```
currency[Greece] = Drachma
```

to remap it to a new value. Unlike a standard array, indices for a map need not be consecutive nor even numeric. Common applications of maps include the following.

- A university's information system relies on some form of a student ID as a key that is mapped to that student's associated record (such as the student's name, address, and course grades) serving as the value.
- The domain-name system (DNS) maps a host name, such as `www.wiley.com`, to an Internet-Protocol (IP) address, such as `208.215.179.146`.
- A social media site typically relies on a (nonnumeric) username as a key that can be efficiently mapped to a particular user's profile.
- A computer graphics system may map a color name, such as `turquoise`, to the triple of numbers that describes the color's RGB (red-green-blue) representation, such as `(64,224,208)`.
- Python uses a dictionary to represent each namespace, mapping an identifying string, such as `pi`, to an associated object, such as `3.14159`.

In this chapter and the next we demonstrate that a map may be implemented so that a search for a key, and its associated value, can be performed very efficiently, thereby supporting fast lookup.

404 Chapter 10. Maps, Hash Tables, and Skip Lists

`M.popitem()` :Remove an arbitrary key-value pair from the map, and return a (k,v) tuple representing the removed pair. If map is empty, raise a `KeyError` .

`M.clear()` :Remove all key-value pairs from the map.

`M.keys()` :Return a set-like view of all keys of M.

`M.values()` :Return a set-like view of all values of M.

`M.items()` :Return a set-like view of (k,v) tuples for all entries of M.

`M.update(M2)` :Assign `M[k] = v` for every (k,v) pair in map M2.

`M == M2` :Return True if maps M and M2 have identical key-value associations.

`M != M2` :Return True if maps M and M2 do not have identical key-value associations.

Example 10.1: In the following, we show the effect of a series of operations on an initially empty map storing items with integer keys and single-character values.

We use the literal syntax for Python's dictclass to describe the map contents.

Operation

Return Value

Map

`len(M)`

0

`{}`

`M[`

`K`

`]=2`

`.`

`{`

`K`

`:2}`

`M[`

`B`

`]=4`

`.`

`{`

`K`

`:2 ,`

`B`

`:4}`

`M[`

`U`

`]=2`

`.`

`{`

`K`

`:2 ,`

`B`

`:4 ,`

`U`

`:2}`

`M[`

10.1.3 Python's MutableMapping Abstract Base Class

Section 2.4.3 provides an introduction to the concept of an abstract base class and the role of such classes in Python's collections module. Methods that are declared to be abstract in such a base class must be implemented by concrete subclasses. However, an abstract base class may provide concrete implementation of other methods that depend upon use of the presumed abstract methods. (This is an example of the template method design pattern.)

The collections module provides two abstract base classes that are relevant to our current discussion: the Mapping and MutableMapping classes. The Mapping class includes all nonmutating methods supported by Python's dict class, while the MutableMapping class extends that to include the mutating methods. What we define as the map ADT in Section 10.1.1 is akin to the MutableMapping abstract base class in Python's collections module.

The significance of these abstract base classes is that they provide a framework to assist in creating a user-defined map class. In particular, the MutableMapping class provides concrete implementations for all behaviors other than the `__delitem__` outlined in Section 10.1.1:

```
getitem
,
setitem
,
delitem
,
len
, and
iter
```

. As we implement the map abstraction with various data structures, as

long as we provide the `__delitem__` core behaviors, we can inherit all other derived behaviors by simply declaring MutableMapping.

To better understand the MutableMapping class, we provide a few examples of

how concrete behaviors can be derived from the `__delitem__` core abstractions. For example, the `contains`

method, supporting the syntax `key in M`, could be implemented by making a guarded attempt to retrieve `self[k]` to determine if the key exists.

```
def
contains
(self, k) :
try:
self[k] # access via
getitem
(ignore result)
return True
except KeyError:
return False # attempt failed
```

A similar approach might be used to provide the logic of the `setdefault` method.

```
def setdefault( self, k, d ) :
```

```
try:
return self [k] # if
getitem
succeeds, return value
```

408 Chapter 10. Maps, Hash Tables, and Skip Lists

1classMapBase(MutableMapping):

2...Our own abstract base class that includes a nonpublic

Item class...

3

4#----- nested

Item class -----

5class

Item:

6 ...Lightweight composite to store key-value pairs as map items...

7

slots

=

_key

,

_value

89 def

init

(self,k ,v) :

10 self.

key = k

11 self.

value = v

1213 def

eq

(self,o t h e r) :

14 return self .

key == other.

key # compare items based on their keys

1516 def

ne

(self,o t h e r) :

17 return not (self== other) #o p p o s i t e o f

eq

18

19 def

lt

(self,o t h e r) :

20 return self .

key<other.

key # compare items based on their keys

Code Fragment 10.2: Extending the MutableMapping abstract base class to provide a nonpublic

Item class for use in our various map implementations.

10.1.5 Simple Unsorted Map Implementation

We demonstrate the use of the MapBase class with a very simple concrete implementation of the map ADT. Code Fragment 10.3 presents an UnsortedTableMap class that relies on storing key-value pairs in arbitrary order within a Python list.

An empty table is initialized as self.

table within the constructor for our map.

426 Chapter 10. Maps, Hash Tables, and Skip Lists

```
26def
bucket
    getitem( self,j ,k ) :
27 found, s = self.
    .nd
    slot(j, k)
28 if not found:
29 raiseKeyError(
    Key Error:
    +r e p r ( k ) ) # no match found
30 return self .
table[s].
    value
31
32def
bucket
    setitem( self,j ,k ,v ) :
33 found, s = self.
    .nd
    slot(j, k)
34 if not found:
35 self.
    table[s] = self.
    Item(k,v) # insert new item
36 self.
    n+ =1 # size has increased
37 else:
38 self.
    table[s].
        value = v # overwrite existing
3940def
bucket
    delitem( self,j ,k ) :
41 found, s = self.
    .nd
    slot(j, k)
42 if not found:
43 raiseKeyError(
    Key Error:
    +r e p r ( k ) ) # no match found
44 self.
    table[s] = ProbeHashMap.
    AVAIL # mark as vacated
4546def
iter
(self):
47 forjinrange(len( self.
table)): # scan entire table
48 if not self .
```

10.3. Sorted Maps 427

10.3 Sorted Maps

The traditional map ADT allows a user to look up the value associated with a given key, but the search for that key is a form known as an exact search .

For example, computer systems often maintain information about events that have occurred (such as financial transactions), organizing such events based upon what are known as time stamps . If we can assume that time stamps are unique for a particular system, then we might organize a map with a time stamp serving as the key, and a record about the event which they occur, or to search for which event occurred closest to a particular time.

In fact, the fast performance of hash-based implementations of the map ADT relies on the intentional scattering of keys.

In this section, we introduce an extension known as the sorted map ADT that includes all behaviors of the standard map, plus the following:

M.min()

:Return the (key,value) pair with minimum key (or None, if map is empty).

M.max()

:Return the (key,value) pair with maximum key (or None, if map is empty).

M.lt(k)

:Return the (key,value) pair with the greatest key that is strictly less than k (or None, if no such key exists).

M.le(k)

:Return the (key,value) pair with the greatest key that is less than or equal to k (or None, if no such key exists).

M.gt(k)

:Return the (key,value) pair with the least key that is strictly greater than k (or None, if no such key exists).

M.ge(k)

:Return the (key,value) pair with the least key that is greater than or equal to k (or None, if no such key exists).

M.range(start, stop)

:Iterate all (key,value) pairs with start ≤ key < stop.

If start is None, iteration begins with minimum key; if

stop is None, iteration concludes with maximum key.

iter(M) :Iterate all keys of the map according to their natural order, from smallest to largest.

reversed(M) :Iterate all keys of the map in reverse order; in Python, this is implemented with the

reversed
method.

430 Chapter 10. Maps, Hash Tables, and Skip Lists

```
1class SortedTableMap(MapBase):
```

```
2...Map implementation using a sorted table...
```

```
3
```

```
4#----- nonpublic behaviors -----
```

```
5def
```

```
·nd
```

```
index(self, k, low, high):
```

```
6...Return index of the leftmost item with key greater than or equal to k.
```

```
78 Return high + 1 if no such item qualifies.
```

```
9
```

```
10 That is, j will be returned such that:
```

```
11 all items of slice table[low:j] have key < k
```

```
12 all items of slice table[j:high+1] have key >= k
```

```
13 ...
```

```
14 if high < low:
```

```
15 return high + 1 # no element qualifies
```

```
16 else:
```

```
17 mid = (low + high) // 2
```

```
18 if k == self.
```

```
table[mid].
```

```
key:
```

```
19 return mid # found exact match
```

```
20 elif k < self.
```

```
table[mid].
```

```
key:
```

```
21 return self .
```

```
·nd
```

```
index(k, low, mid - 1) # Note: may return mid
```

```
22 else:
```

```
23 return self .
```

```
·nd
```

```
index(k, mid + 1, high) # answer is right of mid
```

```
2425 #----- public behaviors -----
```

```
26def
```

```
init
```

```
(self):
```

```
27...Create an empty map...
```

```
28 self.
```

```
table = []
```

```
2930def
```

```
len
```

```
(self):
```

```
31...Return number of items in the map...
```

```
32 return len(self.
```

```
table)
```

```
3334def
```

```
getitem
```

```
(self, k):
```

```
35...Return value associated with key k (raise KeyError if not found)...
```

432 Chapter 10. Maps, Hash Tables, and Skip Lists

```
78def·nd
ge(self,k ) :
79 ...Return (key,value) pair with least key greater than or equal to k...
80 j=self.
·nd
index(k, 0, len( self.
table) ·1) #j
sk e y>=k
81 ifj<len(self.
table):
82 return (self.
table[j].
key,self.
table[j].
value)
83 else:
84 return None
85
86def·nd
lt(self,k ) :
87 ...Return (key,value) pair with greatest key strictly less than k...
88 j=self.
·nd
index(k, 0, len( self.
table) ·1) #j
sk e y>=k
89 ifj>0:
90 return (self.
table[j ·1].
key,self.
table[j ·1].
value) # Note use of j-1
91 else:
92 return None
9394def·nd
gt(self,k ) :
95 ...Return (key,value) pair with least key strictly greater than k...
96 j=self.
·nd
index(k, 0, len( self.
table) ·1) #j
sk e y>=k
97 ifj<len(self.
table) and self .
table[j].
key == k:
98 j+ =1 # advanced past match
99 ifj<len(self.
table):
```

10.3.2 Two Applications of Sorted Maps

In this section, we explore applications in which there is particular advantage to using a sorted map rather than a traditional (unsorted) map. To apply a sorted map, keys must come from a domain that is totally ordered. Furthermore, to take advantage of the inexact or range searches afforded by a sorted map, there should be some reason why nearby keys have relevance to a search.

Flight Databases

There are several Web sites on the Internet that allow users to perform queries on flight databases to find flights between two cities and to buy a ticket. To make a query, a user specifies origin and destination cities, a departure date, and a departure time. To support such queries, we can model the flight database as a map, where keys are tuples of these four parameters. That is, a key is a tuple

$k = (\text{origin}, \text{destination}, \text{date}, \text{time})$.

Additional information about a flight, such as the flight number, the number of seats still available in first (F) and coach (C) class, can be stored in the value object.

Finding a requested flight is not simply a matter of finding an exact match

for a requested query. Although a user typically wants to exactly match the origin and destination cities, he or she may have flexibility for the departure date, and certainly will have some flexibility for the departure time.

We can handle such a query by ordering our keys lexicographically. Then, an efficient implementation for a sorted map would be a good way to satisfy users' queries. For instance, given a user query q , we can find the first flight between the desired cities, having a departure date and time matching the desired query or later. Because the keys are ordered lexicographically, we can use

`range(k1, k2)` to find all flights within a given range of times. For example, if $k1 = (\text{ORD}, \text{PVD}, \text{05May}, \text{09:30})$, and $k2 = (\text{ORD}, \text{PVD}, \text{05May}, \text{20:00})$, a respective call to

`range(k1, k2)` might result in the following sequence of key-value pairs:

`range(k1, k2)` might result in the following sequence of key-value pairs:

`range(k1, k2)` might result in the following sequence of key-value pairs:

`range(k1, k2)` might result in the following sequence of key-value pairs:

`range(k1, k2)` might result in the following sequence of key-value pairs:

`range(k1, k2)` might result in the following sequence of key-value pairs:

`range(k1, k2)` might result in the following sequence of key-value pairs:

`range(k1, k2)` might result in the following sequence of key-value pairs:

`range(k1, k2)` might result in the following sequence of key-value pairs:

`range(k1, k2)` might result in the following sequence of key-value pairs:

436 Chapter 10. Maps, Hash Tables, and Skip Lists

Maintaining a Maxima Set with a Sorted Map

We can store the set of maxima pairs in a sorted map, *M*, so that the cost is the key -eld and performance (speed) is the value -eld. We can then implement operations *add*(*c*,*p*), which adds a new cost-performance pair (*c*,*p*), and *best*(*c*), which returns the best pair with cost at most *c*, as shown in Code Fragment 10.11.

```
1 class CostPerformanceDatabase:
2     ... Maintain a database of maximal (cost, performance) pairs. ...
3
4     def
5         init
6         (self):
7             ... Create an empty database. ...
8             self.
9             M = SortedTableMap( ) # or a more efficient sorted map
10
11     def best(self, c) :
12         ... Return (cost, performance) pair with largest cost not exceeding c.
13         Return None if there is no such pair.
14
15     def add(self, c, p) :
16         ... Add new entry with cost c and performance p. ...
17         # determine if (c, p) is dominated by an existing pair
18         other = self.
19         M.nd
20         le(c) # other is at least as cheap as c
21         if other is not None and other[1] >= p : # if its performance is as good,
22             return # (c, p) is dominated, so ignore
23         self.
24         M[c] = p # else, add (c, p) to database
25         # and now remove any pairs that are dominated by (c, p)
26         other = self.
27         M.nd
28         gt(c) # other more expensive than c
29         while other is not None and other[1] <= p :
30             del self.
31             M[other[0]]
32             other = self.
33         M.nd
34         gt(c)
```

Code Fragment 10.11: An implementation of a class maintaining a set of maxima cost-performance pairs using a sorted map.

Unfortunately, if we implement *M* using the *SortedTableMap*, the *add* behavior has $O(n)$ worst-case running time. If, on the other hand, we implement *M* using a skip list, which we next describe, we can perform *best*(*c*) queries in $O(\log n)$ expected time and *add*(*c*,*p*) updates in $O((1+r)\log n)$ expected time, where *r* is the number of points removed.

438 Chapter 10. Maps, Hash Tables, and Skip Lists

and place that item in S_{i+1} if the coin comes up heads. Thus, we expect S_1 to have about $n/2$ items, S_2 to have about $n/4$ items, and, in general, S_i to have about $n/2^i$ items. In other words, we expect the height of S to be about $\log n$. The halving of the number of items from one list to the next is not enforced as an explicit property of skip lists, however. Instead, randomization is used.

Functions that generate numbers that can be viewed as random numbers are built into most modern computers, because they are used extensively in computer games, cryptography, and computer simulations. Random number generators generate random-like numbers, starting with an initial seed. (See discussion of random module in Section 1.11.1.) Other methods use hardware devices to extract true random numbers from nature. In any case, we will assume that our computer has a random number generator. The main advantage of using randomization in data structure and algorithm design is that the structures and functions that result are usually simple and efficient.

The skip list has the same logarithmic time bounds for searching as is achieved by the binary search algorithm, yet it extends that performance to update methods when inserting or deleting items. In a skip list, while binary search has a worst-case bound with a sorted table.

A skip list makes random choices in arranging its structure in such a way that search and update times are $O(\log n)$ on average, where n is the number of items in the map. Interestingly, the notion of average time complexity used here does not depend on the probability distribution used to help decide where to place the new item. The running time is averaged over all possible outcomes of the random numbers used when inserting entries.

Using the position abstraction used for lists and trees, we view a skip list as a two-dimensional collection of positions arranged horizontally into levels and vertically into towers. Each level is a list S_i

and each tower contains positions storing

the same item across consecutive lists. The positions in a skip list can be traversed using the following operations:

$\text{next}(p)$: Return the position following p on the same level.

$\text{prev}(p)$: Return the position preceding p on the same level.

$\text{below}(p)$: Return the position below p in the same tower.

$\text{above}(p)$: Return the position above p in the same tower.

We conventionally assume that the above operations return `None` if the position

requested does not exist. Without going into the details, we note that we can easily implement a skip list by means of a doubly linked list structure. The structure is essentially a collection of doubly linked lists aligned at towers, which are also doubly linked lists.

Algorithm SkipSearch(k) :

Input: A search key k

Output: Position p in the bottom list S₀ with the largest key such that key(p) ≤ k

p ← start {begin at start position }

while below(p) ≠ None do

p ← below(p) {drop down }

while k < key(next(p)) do

p ← next(p) {scan forward }

return p.

Code Fragment 10.12: Algorithm to search a skip list S for key k.

As it turns out, the expected running time of algorithm SkipSearch on a skip list with n entries is $O(\log n)$. We postpone the justification of this fact, however, until after we discuss the implementation of the update methods for skip lists. Navigation starting at the position identified by SkipSearch(k) can be easily used to provide the additional forms of searches in the sorted map ADT (e.g., find range).

Insertion into a Skip List

The execution of the map operation $M[k] = v$ begins with a call to SkipSearch(k).

This gives us the position p of the bottom-level item with the largest key less than or equal to k (note that p may hold the special item with key ∞). If key(p) = k, the

associated value is overwritten with v. Otherwise, we need to create a new tower for item (k, v). We insert (k, v) immediately after position p within S₀. After inserting

the new item at the bottom level, we use randomization to decide the height of the tower for the new item. We flip a coin, and if the coin comes up tails, then we stop here. Else (the coin comes up heads), we backtrack to the previous (next higher)

level and insert (k, v) in this level at the appropriate position. We again flip a coin;

if it comes up heads, we go to the next higher level and repeat. Thus, we continue to insert the new item (k, v) in list

We link together all the references to the new item (k, v) created in this process to

create its tower. A coin flip can be simulated with Python's built-in pseudo-random number generator from the random module, which returns 0 or 1, each with probability 1/2.

We give the insertion algorithm for a skip list S in Code Fragment 10.13 and

we illustrate it in Figure 10.12. The algorithm uses an insertAfterAbove(p, q, (k, v))

method that inserts a position storing the item (k, v) after position p (on the same

level as p) and above position q, returning the new position r (and setting internal references so that next, prev, above, and below methods will work correctly for p,

q, and r). The expected running time of the insertion algorithm on a skip list with n entries is $O(\log n)$, which we show in Section 10.4.2.

Removal in a Skip List

Like the search and insertion algorithms, the removal algorithm for a skip list is quite simple. In fact, it is even easier than the insertion algorithm. That is, to perform the map operation $\text{del } M[k]$ we begin by executing method `SkipSearch(k)`.

If the position `p` stores an entry with key different from k , we raise a `KeyError`.

Otherwise, we remove `p` and all the positions above `p`, which are easily accessed by using above operations to climb up the tower of this entry in `S` starting at position `p`.

While removing levels of the tower, we reestablish links between the horizontal neighbors of each removed position. The removal algorithm is illustrated in Figure 10.13 and a detailed de-

$O(\log n)$ expected running time.

Before we give this analysis, however, there are some minor improvements to

the skip-list data structure we would like to discuss. First, we do not actually need to store references to values at

more efficiently represent a tower as a single object, storing the key-value pair,

and maintaining `jprevious` references and `jnext` references if the tower reaches

level `S`

j. Second, for the horizontal axes, it is possible to keep the list singly linked,

storing only the next references. We can perform insertions and removals in strictly a top-down, scan-forward fashion.

Exercise C-10.44. Neither of these optimizations improve the asymptotic performance of skip lists by more than a constant factor, but these improvements can, nevertheless, be meaningful in practice.

search trees, which are discussed in Chapter 11.

31 S5

S4

S3

S2

S1---

-- 1212 --

1717 25

25 31

3142

55 5055+

+..+

+

+..

+....

17

38

38 39 424242

44

445555+

17

17

20 2525

S0

Figure 10.13: Removal of the entry with key 25 from the skip list of Figure 10.12.

The positions visited after the search for the position of `S0` holding the entry are

highlighted. The positions removed are drawn with dashed lines.

Bounding the Height of a Skip List

Because the insertion step involves randomization, a more accurate analysis of skip lists involves a bit of probability. At first, this might seem like a major undertaking, for a complete and thorough probabilistic analysis could require deep mathematics (and, indeed, there are several ways to understand the expected asymptotic behavior of skip lists. The informal and intuitive probabilistic analysis we give below uses only basic concepts of probability theory. Let us begin by determining the expected value of the height of a skip list S with n entries (assuming that we do not terminate insertions early). The probability that a given entry has a tower of height $i+1$ is equal to the probability of getting i consecutive heads when flipping a coin, that is, this probability is $1/2^i$.

Hence, the probability that level i has at least one position is at most

$$P_i \leq n/2^i,$$

for the probability that any one of n different events occurs is at most the sum of the probabilities that each occurs.

The probability that the height of S is larger than i is equal to the probability that level i has at least one position, that is, it is no more than P_i . This means that h is larger than, say, $3 \log n$ with probability at most

$$P_{3 \log n} \leq n/2^{3 \log n}$$

$$= n/2^{\log n^3}$$

$$= n/n^3$$

$$= 1/n^2.$$

For example, if $n=1000$, this probability is a one-in-a-million long shot. More generally, given a constant $c > 1$, h is larger than $c \log n$ with probability at most $1/n^{c-1}$. That is, the probability that h is smaller than $c \log n$ is at least $1 - 1/n^{c-1}$.

Thus, with high probability, the height of S is $O(\log n)$.

Analyzing Search Time in a Skip List

Next, consider the running time of a search in skip list S , and recall that such a search involves two nested while loops. The inner loop performs a scan forward on all elements of S as long as the next key is no greater than the search key k , and the outer loop drops down to the next level and repeats the scan forward iteration. Since the height of S is $O(\log n)$ with high probability, the search time is $O(\log n)$ with high probability.

10.5.2 Python's MutableSet Abstract Base Class

To aid in the creation of user-defined set classes, Python's `collections` module provides a `MutableSet` abstract base class (just as it provides the `MutableMapping` abstract base class discussed in Section 10.1.3). The `MutableSet` base class provides concrete implementations for all methods described in Section 10.5.1, except for the core behaviors (`add`, `discard` , `contains`

, `len`, `__and__`, `__iter__`) that must

be implemented by any concrete subclass. This design is an example of what is known as the template method pattern , as the concrete methods of the `MutableSet` class rely on the presumed abstract methods that will subsequently be provided by a subclass.

For the purpose of illustration, we examine algorithms for implementing several of the derived methods of the `MutableSet` base class. For example, to determine if one set is a proper subset of another, we must verify two conditions: a proper subset must have size strictly smaller than that of its superset, and each element of a subset must be contained in the superset. An implementation of the corresponding

method based on this logic is given in Code Fragment 10.14.

```
def
It
(self, other) : # supports syntax S < T
...Return true if this set is a proper subset of other.
...if len(self) >= len(other):
return False # proper subset must have strictly smaller size
for element in self :
    if element not in other:
return False # not a subset since element missing from other
return True # success; all conditions are met
```

Code Fragment 10.14: A possible implementation of the `MutableSet`.

It
method, which tests if one set is a proper subset of another.

As another example, we consider the computation of the union of two sets.

The set ADT includes two forms for computing a union. The syntax `S|T` should produce a new set that has contents equal to the union of existing sets `S` and `T`. This operation is implemented through the special method

`or`
in Python. Another

syntax, `S|=T` is used to update existing set `S` to become the union of itself and set `T`. Therefore, all elements of `T` that are not already contained in `S` should be added to `S`. We note that this 'in-place' operation may be implemented more efficiently than if we were to rely on the 'first' form, using the syntax `S=S|T`, in

which identifier `S` is reassigned to a new set instance that represents the union. For

convenience, Python's built-in set class supports named version of these behaviors, with `S.union(T)` equivalent to those named versions are not formally provided by the `MutableSet` abstract base class).

Sets

Although sets and maps have very different public interfaces, they are really quite similar. A set is simply a map in which keys do not have associated values. Any data structure used to implement a set is storing set elements as keys, and using None as an irrelevant value, but such an implementation is unnecessarily wasteful. An efficient set implementation should abandon the

Item composite that we use in our MapBase class and instead store set elements directly in a data structure.

Multisets

The same element may occur several times in a multiset. All of the data structures we have seen can be reimplemented in which the map key is a (distinct) element of the multiset, and the associated value is a count of the number of occurrences.

Python's standard collections module includes a definition for a class named Counter that is in essence a multiset. Formally, the Counter class is a subclass of dict, with the expectation that values are integers, and with additional functionality like a most

common(n) method that returns a list of the n most common elements.

The standard

iter

reports each element only once (since those are formally the keys of the dictionary). There is another method named elements() that iterates through the multiset with each element being repeated according to its count.

Multimaps

Although there is no multimap in Python's standard libraries, a common implementation approach is to use a standard dictionary. This class uses the standard dict class as the map, and a list of values as a composite value in the dictionary. We have designed the class so that a different map implementation can easily be substituted by overriding the MapType attribute at line 3.

456 Chapter 10. Maps, Hash Tables, and Skip Lists

C-10.39 Although keys in a map are distinct, the binary search algorithm can be applied in a more general setting in which an array stores possibly duplicative elements in nondecreasing order. Consider the goal of identifying the index of the leftmost element with key g .

Does the

`find`

index method as given in Code Fragment 10.8 guarantee

such a result? Does the

`find`

index method as given in Exercise R-10.21

guarantee such a result? Justify your answers.

C-10.40 Suppose we are given two sorted search tables S and T , each with n entries (with S and T being implemented with arrays). Describe an $O(\log^2 n)$ -time algorithm for finding the k th smallest key in the union of the keys from S and T (assuming no duplicates).

C-10.41 Give an $O(\log n)$ -time solution for the previous problem.

C-10.42 Suppose that each row of an $n \times n$ array A consists of 1-s and 0-s such that, in any row of A , all the 1-s come before any 0-s in that row. Assuming A is already in memory, describe a method running in $O(n \log n)$ time (not $O(n^2)$ time!) for counting the number of 1-s in A .

C-10.43 Given a collection C of cost-performance pairs (c, p) , describe an algorithm for finding the maxima pairs of C in $O(n \log n)$ time.

C-10.44 Show that the methods `above(p)` and `prev(p)` are not actually needed to efficiently implement a map using a skip list. That is, we can implement insertions and deletions in a skip list using only `find`. In the insertion algorithm, first repeatedly flip the coin to determine the level where you should start inserting the new entry.

C-10.45 Describe how to modify a skip-list representation so that index-based operations, such as retrieving the item at index j , can be performed in $O(\log n)$ expected time.

C-10.46 For sets S and T , the syntax $S \cdot T$ returns a new set that is the symmetric difference, that is, a set of elements that are in precisely one of S or T . This syntax is supported by the special

`xor`

method. Provide an

implementation of that method in the context of the `MutableSet` abstract base class, relying only on the `remove` primary abstract methods of that class.

C-10.47 In the context of the `MutableSet` abstract base class, describe a concrete implementation of the

`intersect`

method, which supports the syntax $S \& T$

for computing the intersection of two existing sets.

C-10.48 An inverted file is a critical data structure for implementing a search engine or the index of a book. Given a document D , which can be viewed as an unordered, numbered list of words, an inverted file is an ordered list of words, L , such that, for each word w in L , we store the indices of the places in D where w appears. Design an efficient algorithm for constructing L from D .

Chapter Notes

Hashing is a well-studied technique. The reader interested in further study is encouraged to explore the book by Knuth [65], as well as the book by Vitter and Chen [100]. Skip lists were introduced by Pugh [86]. Our analysis of skip lists is a simplification of a presentation given by Motwani and Raghavan [80]. For a more in-depth analysis of skip lists, please see the various research papers on skip lists that have appeared in the data structures literature [59, 81, 84]. Exercise C-10.36 was contributed by James Lee.

another occurrence. The efficiency of the Boyer-Moore algorithm relies on creating a lookup table that quickly determines where a mismatched character occurs elsewhere in the pattern. In particular, we define a function `last(c)` as follows: If `c` is in `P`, `last(c)` is the index of the last (rightmost) occurrence of `c` in `P`. Otherwise, we conventionally define `last(c) = -1`.

If we assume that the alphabet is of fixed, finite size, and that characters can be converted to indices of an array (for example, by using their character code), the `last` function can be easily implemented as a lookup table with worst-case $O(1)$ -time access to the value `last(c)`. However, the table would have length equal to the size of the alphabet (rather than the size of the pattern), and time would be required to initialize the entire table.

We prefer to use a **hash** table to represent the `last` function, with only those characters from the pattern occurring in the structure. The space usage for this approach is proportional to the number of characters in the pattern, and thus $O(m)$. The expected lookup time remains independent of the problem (although the worst-case bound is $O(m)$). Our complete implementation of the Boyer-Moore pattern-matching algorithm is given in Code Fragment 13.2.

```

1 def find
2   boyer
3   moore(T, P):
4     ...Return the lowest index of T at which substring P begins (or else -1)...
5     3n, m = len(T), len(P) # introduce convenient notations
6     4if m == 0 : return 0 # trivial search for empty string
7     5last = {} # build 'last' dictionary
8     6for k in range(m):
9       7last[P[k]] = k # later occurrence overwrites
10    8# align end of pattern at index m-1 of text
11    9i = m - 1 # an index into T
12    10 k = m - 1 # an index into P
13    11while i < n:
14      12 if T[i] == P[k]: # a matching character
15        13 if k == 0 :
16          14 return i # pattern begins at index i of text
17        15 else:
18          16 i -= 1 # examine previous character
19          17 k -= 1 # offset to hTandP
20        18 else:
21          19 j = last.get(T[i], -1) # last(T[i]) is -1 if not found
22          20 i += m - min(k, j + 1) # case analysis for jump step
23          21 k = m - 1 # restart at end of pattern
24    22return -1

```

Code Fragment 13.2: An implementation of the Boyer-Moore algorithm.

13.2. Pattern-Matching Algorithms 589

The correctness of the Boyer-Moore pattern-matching algorithm follows from the fact that each time the method makes a shift, it is guaranteed not to skip over any possible matches. For last(c) is the location of the last occurrence of c in P.

In Figure 13.4, we illustrate the execution of the Boyer-Moore pattern-matching algorithm on an input string similar to

```
c
abc d
last (c)
453 1
a c d a b a a cb aabc Text:
Pattern: b aaa b c1
baaa b c2 3 4
baaa b c5
baaa b c6baaa b c7baaa b c8 9 10 11 12 13b b a a b a
```

Figure 13.4: An illustration of the Boyer-Moore pattern-matching algorithm, including a summary of the last(c) function. The algorithm performs 13 character comparisons, which are indicated with numerical labels.

Performance

If using a traditional lookup table, the worst-case running time of the Boyer-Moore algorithm is $O(nm + |\Sigma|)$. Namely, $O(m + |\Sigma|)$, and the actual search for the pattern takes $O(nm)$ time in the worst case, the same as the brute-force algorithm. (With a hash table, the dependence on $|\Sigma|$ is removed.) An example of a text-pattern pair that achieves the worst case is

```
T = n/bracehtipdownleft
/bracehtipupright/bracehtipupleft
/bracehtipdownrightaaaaa ...a
P = bm-1/bracehtipdownleft
/bracehtipupright/bracehtipupleft
/bracehtipdownrightaa...a
```

The worst-case performance, however, is unlikely to be achieved for English text, for, in that case, the Boyer-Moore comparisons done per character is 0.24 for a 26-character pattern string.

We have actually presented a simplified version of the Boyer-Moore algorithm.

The original algorithm achieves running time $O(n + m + |\Sigma|)$ by using an alternative shift heuristic to the partially matched text string, whenever it shifts the pattern more than the character-jump heuristic. This alternative shift heuristic is based on applying the main idea from the Knuth-Morris-Pratt pattern-matching algorithm, which we discuss next.

14.2.5 Python Implementation

In this section, we provide an implementation of the Graph ADT. Our implementa-

tion will support directed or undirected graphs, but for ease of explanation, we first describe it in the context of an undirected graph.

We use a variant of the adjacency map representation. For each vertex v , we

use a Python dictionary to represent the secondary incidence map $I(v)$. However,

we do not explicitly maintain lists V and E , as originally described in the edge list

representation. The list V is replaced by a top-level dictionary D that maps each

vertex v to its incidence map $I(v)$; note that we can iterate through all vertices by

generating the set of keys for dictionary D . By using such a dictionary D to map

vertices to the secondary incidence maps, we need not maintain references to those incidence maps as part of the

$O(1)$ expected time. This greatly simplifies our implementation. However, a

consequence of our design is that some of the worst-case running time bounds for the graph ADT operations, given

by the edge list representation, are no longer valid. Instead of maintaining list E , we are content with taking the union of the edges found in the

various incidence maps; technically, this runs in $O(n+m)$ time rather than strictly

$O(m)$ time, as the dictionary D has n keys, even if some incidence maps are empty.

Our implementation of the graph ADT is given in Code Fragments 14.1 through

14.3. Classes `Vertex` and `Edge`, given in Code Fragment 14.1, are rather simple,

and can be nested within the more complex `Graph` class. Note that we define the

`hash`

method for both `Vertex` and `Edge` so that those instances can be used as

keys in Python's `hash`-based sets and dictionaries. The rest of the `Graph` class is

given in Code Fragments 14.2 and 14.3. Graphs are undirected by default, but can be declared as directed with a

parameter. Internally, we manage the directed case by having two different top-level dictionary

instances,

`outgoing` and

`incoming`, such that

`outgoing[v]` maps to another

dictionary representing $\text{out}(v)$, and

`incoming[v]` maps to a representation of $\text{in}(v)$.

In order to unify our treatment of directed and undirected graphs, we continue to use the

`outgoing` and

`incoming` identifiers in the undirected case, yet as aliases

to the same dictionary. For convenience, we define a utility named `is`

`directed` to

allow us to distinguish between the two cases.

For methods `degree` and `incident`

`edges`, which each accept an optional param-

eter to differentiate between the outgoing and incoming orientations, we choose the appropriate map before processing

vertex v ; we always initial-

ize

`outgoing[v]` to an empty dictionary for new vertex v . In the directed case, we

independently initialize

`incoming[v]` as well. For the undirected case, that step is

unnecessary as

`outgoing` and

`incoming` are aliases. We leave the implementations

of methods `remove`

`vertex` and `remove`

`edge` as exercises (C-14.37 and C-14.38).

Index

#character, 3

·operator, 14, 75

%operator, 13·14, 75, 242

&operator, 14, 75

operator, 13, 14, 75

operator, 75

=operator, 75

+operator, 13, 14, 75

+=operator, 16, 75

·operator, 13, 75

·=operator, 75

/operator, 13, 75

//operator, 13·14, 75

<operator, 13, 15, 75, 76

<<operator, 14, 75, 413

<=operator, 13, 15, 75, 76

=operator, 4

==operator, 12, 15, 75, 76

>operator, 13, 15, 75, 76

>=operator, 13, 15, 75

>>operator, 14, 75, 413

·operator, 14, 75, 412

abs

,7 5

add

, 74·76

and

,7 5

bool

, 74·76

call

,7 5

contains

,75, 76, 95, 203, 403

delitem

, 75, 403, 460

eq

, 75, 76

·oat

,7 5

·oordiv

,7 5

ge

,7 5

getitem

, 75, 79, 80, 93, 95, 211,

212, 403, 460

gt

,7 5

Index 745

open function, 29, 31

operand stack, 704

operators, 12-17

arithmetic, 13

bitwise, 14

chaining, 17comparisons, 13logical, 12

overloading, 74

precedence, 17

operator, 12

ordfunction, 29

order statistic, 571

OrderedDict class, 457

Orlin, James, 696

osmodule, 49, 159, 182, 357

out-degree, 621

outgoing edge, 621

over-ow, 506

override, 82, 100

p-norm, 53

packing a sequence, 44

palindrome, 181, 615

parameter, 24-28

actual, 24

default value, 26formal, 24

keyword, 27

positional, 27

parent class, 82

parent node, 301

parenthetic string representation, 339

partial order, 16

partition, 679, 681-684

pass statement, 38, 478

path, 302, 623

compression, 684

directed, 623length, 356, 660simple, 623

pattern matching, 208, 584-593

Boyer-Moore algorithm, 586-589

brute force, 584-585Knuth-Morris-Pratt algorithm,
590-593

Patterson, David, 719Perkovic, Ljubomir, 55

permutation, 150Peters, Tim, 568

Phillips, Dusty, 108polymorphism, 26, 77, 93polynomial function, 119, 146polynomial **hash** code, 413

portability, 58

position, 279-281, 305, 438

positional list, 277-285

abstract data type, 279-281

PositionalList class, 281-285, 287, 628

positional parameter, 27post-x notation, 252, 253, 359

postorder tree traversal, 329

644 Chapter 14. Graph Algorithms

14.3.2 DFS Implementation and Extensions

We begin by providing a Python implementation of the basic depth-first search algorithm, originally described with pseudo-code in Code Fragment 14.4. Our DFS function is presented in Code Fragment 14.5.

```
1defDFS(g, u, discovered):
2    ...Perform DFS of the undiscovered portion of Graph g starting at Vertex u.
3
4    discovered is a dictionary mapping each vertex to the edge that was used to
5    discover it during the DFS. (u should be 'discovered' prior to the call.)
6    Newly discovered vertices will be added to the dictionary as a result.
7    ...
8    for e in g.incident
9        edges(u): # for every outgoing edge from u
10           v = e.opposite(u)
11           if v not in discovered: # v is an unvisited vertex
12              discovered[v] = e # e is the tree edge that discovered v
13              DFS(g, v, discovered) # recursively explore from v
```

Code Fragment 14.5: Recursive implementation of depth-first search on a graph, starting at a designated vertex u.

In order to track which vertices have been visited, and to build a representation of the resulting DFS tree, our implementation introduces a third parameter, named `discovered`. This parameter is a dictionary mapping each vertex to the tree edge that was used to discover that vertex. As a technicality, we assume that the source vertex `u` is already in the dictionary with a `None` value. Thus, a caller might start the traversal as follows:

```
result = {u:None} # a new dictionary, with u trivially discovered
DFS(g, u, result)
```

The dictionary serves two purposes. Internally, the dictionary provides a mechanism for recognizing visited vertices. The values within the dictionary are the DFS tree edges at the conclusion of the process.

Because the dictionary is `hash`-based, the test, `v not in discovered`, and the record-keeping step, `discovered[v] = e`, run in $O(1)$ expected time, rather than worst-case time. In practice, this is a compromise we are willing to accept, but it does violate the formal analysis. A `discovered` array could be used as indices into an array-based lookup table rather than a `hash`-based map. Alternatively, we could store the tree edge directly as part of the vertex instance.

Chapter Notes 535

Chapter Notes

Some of the data structures discussed in this chapter are extensively covered by Knuth in his *Sorting and Searching* book [65], and by Mehlhorn in [76]. AVL trees are due to Adel-son-Vel'skii and Landis [2], who invented this class of balanced search trees in 1962. Binary search trees, AVL trees, and hashing are described in Knuth's *Sorting and Searching* [65] book. Average-height analyses for binary search trees can be found in the books by Aho, Hopcroft, and Ullman [6] and Cormen, Leiserson, Rivest and Stein [29]. The handbook by Gonnet and Baeza-Yates [44] contains a number of theoretical and experimental comparisons among many trees, which are similar to (2,4) trees. Red-black trees were defined by Bayer [10]. Variations and interesting properties of red-black trees are presented in a paper by Guibas and Sedgewick [48]. The reader interested in learning more about different balanced tree data structures is referred to [76] by Mehlhorn and Tsakalidis [78]. Knuth [65] is excellent additional reading that includes early approaches to balancing trees. Splay trees were invented by Sleator and Tarjan [89] (see also [95]).

736 Bibliography

- [85] R. C. Prim, "Shortest connection networks and some generalizations," *Bell Syst. Tech. J.*, vol. 36, pp. 1389-1401, 1957.
- [86] W. Pugh, "Skip lists: a probabilistic alternative to balanced trees," *Commun. ACM*, vol. 33, no. 6, pp. 668-676, 1990.
- [87] H. Samet, *The Design and Analysis of Spatial Data Structures*. Reading, MA: Addison-Wesley, 1990.
- [88] R. Schaffer and R. Sedgewick, "The analysis of heapsort," *Journal of Algorithms*, vol. 15, no. 1, pp. 76-100, 1993.
- [89] D. D. Sleator and R. E. Tarjan, "Self-adjusting binary search trees," *J. ACM*, vol. 32, no. 3, pp. 652-686, 1985.
- [90] G. A. Stephen, *String Searching Algorithms*. World Scientific Press, 1994.
- [91] M. Summerfeld, *Programming in Python 3: A Complete Introduction to the Python Language*. Addison-Wesley Professional, 2nd ed., 2009.
- [92] R. Tamassia and G. Liotta, "Graph drawing," in *Handbook of Discrete and Computational Geometry* (J. E. Goodman and J. O'Rourke, eds.), ch. 52, pp. 1163-1186, CRC Press LLC, 2nd ed., 2004.
- [93] R. Tarjan and U. Vishkin, "An efficient parallel biconnectivity algorithm," *SIAM J. Comput.*, vol. 14, pp. 862-874, 1985.
- [94] R. E. Tarjan, "Depth-first search and linear graph algorithms," *SIAM J. Comput.*, vol. 1, no. 2, pp. 146-160, 1972.
- [95] R. E. Tarjan, *Data Structures and Network Algorithms*, vol. 4 of CBMS-NSF Regional Conference Series in Applied Mathematics. Philadelphia, PA: Society for Industrial and Applied Mathematics, 1983.
- [96] A. B. Tucker, Jr., *The Computer Science and Engineering Handbook*. CRC Press, 1997.
- [97] J. D. Ullman, *Principles of Database Systems*. Potomac, MD: Computer Science Press, 1983.
- [98] J. van Leeuwen, "Graph algorithms," in *Handbook of Theoretical Computer Science* (J. van Leeuwen, ed.), vol. A. Algorithms and Complexity, pp. 525-632, Amsterdam: Elsevier, 1990.
- [99] J. S. Vitter, "Efficient memory access in large-scale computation," in *Proc. 8th Sympos. Theoret. Aspects Comput. Sci.*, Lecture Notes Comput. Sci., Springer-Verlag, 1991.
- [100] J. S. Vitter and W. C. Chen, *Design and Analysis of Coalesced Hashing*. New York: Oxford University Press, 1987.
- [101] J. S. Vitter and P. Flajolet, "Average-case analysis of algorithms and data structures," in *Algorithms and Complexity* (J. van Leeuwen, ed.), vol. A of *Handbook of Theoretical Computer Science*, pp. 431-524, Amsterdam: Elsevier, 1990.
- [102] S. Warshall, "A theorem on boolean matrices," *Journal of the ACM*, vol. 9, no. 1, pp. 11-12, 1962.
- [103] J. W. J. Williams, "Algorithm 232: Heapsort," *Communications of the ACM*, vol. 7, no. 6, pp. 347-348, 1964.
- [104] D. Wood, *Data Structures, Algorithms, and Performance*. Reading, MA: Addison-Wesley, 1993.
- [105] J. Zelle, *Python Programming: An Introduction to Computer Science*. Franklin, Beedle & Associates Inc., 2nd ed., 2010.

10.5. Sets, Multisets, and Multimaps 449

```
def
  or
  (self, other) : # supports syntax S | T
  ...Return a new set that is the union of two existing sets...
  result = type( self )( ) # create new instance of concrete class
  for e in self :
    result.add(e)
  for e in other:
    result.add(e)
  return result
```

Code Fragment 10.15: An implementation of the MutableSet.

or
method,
which computes the union of two existing sets.

An implementation of the behavior that computes a new set as a union of two others is given in the form of the

or
special method, in Code Fragment 10.15.

An important subtlety in this implementation is the instantiation of the resulting set. Since the MutableSet class is must belong to a concrete subclass. When computing the union of two such concrete instances, the result should presumably be an instance of the same class as the operands. The function `type(self)` returns a reference to the actual class of the instance identified as `self`, and the subsequent parentheses in expression `type(self)()` call the default constructor for that class.

In terms of efficiency, we analyze such set operations while letting n denote the size of S and m denote the size of set T for an operation such as $S|T$. If the concrete sets are implemented with hashing, the expected running time of the implementation in Code Fragment 10.15 is $O(m+n)$, because it loops over both sets, performing constant-time operations in the form of a containment check and a possible insertion into the result.

Our implementation of the in-place version of a union is given in Code Fragment 10.16, in the form of the

or
special method that supports syntax $S|=T$.

Notice that in this case, we do not create a new set instance, instead we modify and return the existing set, after in-place version of the union has expected running time $O(m)$ where m is the size of the second set, because we only have to loop through that second set.

```
def
  or
  (self, other) : # supports syntax S |= T
  ...Modify this set to be the union of itself and another set...
  for e in other:
    self.add(e)
  return self # technical requirement of in-place operator
```

Code Fragment 10.16: An implementation of the MutableSet.

or
method,
which performs an in-place union of one set with another.

14.2.3 Adjacency Map Structure

In the adjacency list structure, we assume that the secondary incidence collections are implemented as unordered linked lists. Such a collection $I(v)$ uses space proportional to $O(\deg(v))$, allows an edge to be added or removed in $O(1)$ time, and allows an iteration of all edges incident to vertex v in $O(\deg(v))$ time. However, the best implementation of `get`

`edge(u,v)` requires $O(\min(\deg(u), \deg(v)))$ time, because we must search through either $I(u)$ or $I(v)$.

We can improve the performance by using a **hash**-based map to implement $I(v)$ for each vertex v . Specifically, we let the opposite endpoint of each incident edge serve as a key in the map, with the edge structure serving as the value. We call such a graph representation an adjacency map. (See Figure 14.6.) The space usage for an adjacency map remains $O(n+m)$, because $I(v)$ uses $O(\deg(v))$ space for each vertex v , as with the adjacency list.

The advantage of the adjacency map, relative to an adjacency list, is that the `get`

`edge(u,v)` method can be implemented in expected $O(1)$ time by searching for vertex u as a key in $I(v)$, or vice versa. This provides a likely improvement over the adjacency list, while retaining the worst-case bound of $O(\min(\deg(u), \deg(v)))$.

In comparing the performance of adjacency map to other representations (see Table 14.1), we find that it essentially achieves optimal running times for all methods, making it an excellent all-purpose choice as a graph representation.

he g
vu
wzf gh
w
huuw v
g e
f ew
vuz
fv
w
zV
(a) (b)

Figure 14.6: (a) An undirected graph G ; (b) a schematic representation of the adjacency map structure for G . Each vertex maintains a secondary map in which neighboring vertices serve as keys, with the connecting edges as associated values. Although not diagrammed as such, we presume that there is a unique `Edge` instance for each edge of the graph, and that it maintains references to its endpoint vertices.