

10.2. **Hash** Tables 415**Hash** Codes in Python

The standard mechanism for computing **hash** codes in Python is a built-in function with signature **hash**(x) that returns an integer value that serves as the **hash** code for object x. However, only immutable data types are deemed hashable in Python. This restriction is meant to ensure that a particular object's **hash** code remains constant during that object's lifespan. This is an important property for an object's use as a key in a **hash** table. A problem could occur if a key were inserted into the hashtable, yet a later search were performed for that key based on a different **hash** code than that which it had when inserted; the wrong bucket would be searched.

Among Python's built-in data types, the immutable `int`, `float`, `str`, `tuple`, and `frozenset` classes produce robust **hash** codes, via the **hash** function, using techniques similar to those discussed earlier in this section. **Hash** codes for character strings are well crafted based on a technique similar to polynomial **hash** codes, except using exclusive-or computations rather than additions. If we repeat the experiment described in Table 10.1 using Python's built-in **hash** codes, we find that only 8 strings out of the set of more than 230,000 collide with another. **Hash** codes

for tuples are computed with a similar technique based upon a combination of the **hash** codes of the individual elements of the tuple. When hashing a `frozenset`, the order of the elements should be irrelevant, and so a natural option is to compute the exclusive-or of the individual **hash** codes without any shifting. If **hash**(x) is called for an instance x of a mutable type, such as a list, a `TypeError` is raised.

Instances of user-defined classes are treated as unhashable by default, with a `TypeError` raised by the **hash** function. However, a function that computes **hash** codes can be implemented in the form of a special method named

`<b>hash</b>`

within

a class. The returned `<b>hash</b>` code should reflect the immutable attributes of an in-stance. It is common to return a `<b>hash</b>` code that is itself based on the computed hash of the combination of such attributes. For example, a Color class that maintains

three numeric red, green, and blue components might implement the method as:

```
def
```

```
<b>hash</b>
```

```
(self):
```

```
return <b>hash</b>( ( self.
```

```
red,self.
```

```
green, self.
```

```
blue) ) # <b>hash</b> combined tuple
```

An important rule to obey is that if a class defines equivalence through

```
eq
```

```
,
```

then any implementation of

```
<b>hash</b>
```

must be consistent, in that if  $x = y$ , then

`<b>hash</b>(x) == <b>hash</b>(y)`. This is important because if two instances are considered

to be equivalent and one is used as a key in a `<b>hash</b>` table, a search for the second instance should result in the discovery of the first. It is therefore important that the hash code for the second match the `<b>hash</b>` code for the first, so that the proper bucket is examined. This rule extends to any well-defined comparisons between objects of different classes. For example, since Python treats the expression  $5 = 5.0$  as

true, it ensures that `<b>hash</b>(5)` and `hash(5.0)` are the same.

412 Chapter 10. Maps, **Hash** Tables, and Skip Lists**Hash** Codes

The first action that a **hash** function performs is to take an arbitrary key  $k$  in our map and compute an integer that is called the **hash** code for  $k$ ; this integer need not be in the range  $[0, N-1]$ , and may even be negative. We desire that the set of **hash** codes assigned to our keys should avoid collisions as much as possible. For if the **hash** codes of our keys cause collisions, then there is no hope for our compression function to avoid them. In this subsection, we begin by discussing the theory of hash codes. Following that, we discuss practical implementations of **hash** codes in Python.

## Treating the Bit Representation as an Integer

To begin, we note that, for any data type  $X$  that is represented using at most as many bits as our integer **hash** codes, we can simply take as a **hash** code for  $x$  an integer interpretation of its bits. For example, the **hash** code for key 314 could simply be 314. The **hash** code for a floating-point number such as 3.14 could be based upon an interpretation of the bits of the floating-point representation as an integer.

For a type whose bit representation is longer than a desired **hash** code, the above scheme is not immediately applicable. For example, Python relies on 32-bit **hash** codes. If a floating-point number uses a 64-bit representation, its bits cannot be viewed directly as a **hash** code. One possibility is to use only the high-order 32 bits (or the low-order 32 bits). This **hash** code, of course, ignores half of the information present in the original key, and if many of the keys in our map only differ in these bits, then they will collide using this simple **hash** code.

A better approach is to combine in some way the high-order and low-order portions of a 64-bit key to form a 32-bit **hash** code, which takes all the original bits into consideration. A simple implementation is to add the two components as 32-bit numbers (ignoring

overflow), or to take the exclusive-or of the two components. These approaches of combining components can be extended to any object  $x$  whose binary representation can be viewed as an  $n$ -tuple  $(x_0, x_1, \dots, x_{n-1})$  of 32-bit integers, for example, by forming a **hash** code for  $x$  as  $\sum_{i=0}^{n-1} x_i$  or as  $x_0 \oplus x_1 \oplus \dots \oplus x_{n-1}$ , where the  $\oplus$  symbol represents the bitwise exclusive-or operation (which is  $\wedge$  in Python).

### Polynomial **Hash** Codes

The summation and exclusive-or **hash** codes, described above, are not good choices for character strings or other variable-length objects that can be viewed as tuples of the form  $(x_0, x_1, \dots, x_{n-1})$ , where the order of the  $x_i$ 's is significant. For example, consider a 16-bit **hash** code for a character string  $s$  that sums the Unicode values of the characters in  $s$ . This **hash** code unfortunately produces lots of unwanted

Rezultat za stranicu 433

## 10.2. **Hash** Tables 411

### 10.2.1 **Hash** Functions

The goal of a **hash** function  $h$ , is to map each key  $k$  to an integer in the range  $[0, N-1]$ , where  $N$  is the capacity of the bucket array for a **hash** table. Equipped with such a **hash** function,  $h$ , the main idea of this approach is to use the **hash** function value,  $h(k)$ , as an index into our bucket array,  $A$ , instead of the key  $k$  (which may not be appropriate for direct use as an index). That is, we store the item  $(k, v)$  in the bucket  $A[h(k)]$ .

If there are two or more keys with the same **hash** value, then two different items will be mapped to the same bucket in  $A$ . In this case, we say that a collision has

occurred. To be sure, there are ways of dealing with collisions, which we will discuss later, but the best strategy is to try to avoid them in the first place. We say that a **hash** function is “good” if it maps the keys in our map so as to sufficiently minimize collisions. For practical reasons, we also would like a **hash** function to be fast and easy to compute.

It is common to view the evaluation of a **hash** function,  $h(k)$ , as consisting of two portions—a **hash** code that maps a key  $k$  to an integer, and a compression function that maps the **hash** code to an integer within a range of indices,  $[0, N-1]$ , for a bucket array. (See Figure 10.5.)

-1 hash code

120 -2... ..

compression function

120N - 1 ...Arbitrary Objects

Figure 10.5: Two parts of a **hash** function: a **hash** code and a compression function.

The advantage of separating the **hash** function into two such components is that the **hash** code portion of that computation is independent of a specific **hash** table size. This allows the development of a general **hash** code for each object that can be used for a **hash** table of any size; only the compression function depends upon the table size. This is particularly convenient, because the underlying bucket array for a **hash** table may be dynamically resized, depending on the number of items currently stored in the map. (See Section 10.2.3.)

Rezultat za stranicu 438

416 Chapter 10. Maps, **Hash** Tables, and Skip Lists

Compression Functions

The **hash** code for a key  $k$  will typically not be suitable for immediate use with a

bucket array, because the integer `hash` code may be negative or may exceed the capacity of the bucket array. Thus, once we have determined an integer `hash` code for a key object `k`, there is still the issue of mapping that integer into the range  $[0, N-1]$ .

This computation, known as a compression function, is the second action performed as part of an overall `hash` function. A good compression function is one that minimizes the number of collisions for a given set of distinct `hash` codes.

### The Division Method

A simple compression function is the division method, which maps an integer `h` to  $h \bmod N$ ,

where  $N$ , the size of the bucket array, is a fixed positive integer. Additionally, if we take  $N$  to be a prime number, then this compression function helps “spread out” the distribution of hashed values. Indeed, if  $N$  is not prime, then there is greater risk that patterns in the distribution of `hash` codes will be repeated in the distribution of hash values, thereby causing collisions. For example, if we insert keys with `hash` codes  $\{200, 205, 210, 215, 220, \dots, 600\}$  into a bucket array of size 100, then each `hash` code will collide with three others. But if we use a bucket array of size 101, then there will be no collisions. If a `hash` function is chosen well, it should ensure that the probability of two different keys getting hashed to the same bucket is  $1/N$ .

Choosing  $N$  to be a prime number is not always enough, however, for if there is a repeated pattern of `hash` codes of the form  $pN + q$  for several different  $p$ 's, then there will still be collisions.

### The MAD Method

A more sophisticated compression function, which helps eliminate repeated patterns in a set of integer keys, is the Multiply-Add-and-Divide (or “MAD”) method.

This method maps an integer `h` to

$$[(a \cdot h + b) \bmod p] \bmod N,$$

where  $N$  is the size of the bucket array,  $p$  is a prime number larger than  $N$ , and  $a$  and  $b$  are integers chosen at random from the interval  $[0, p-1]$ , with  $a > 0$ . This compression function is chosen in order to eliminate repeated patterns in the set of hash codes and get us closer to having a “good” hash function, that is, one such that the probability any two different keys collide is  $1/N$ . This good behavior would be the same as we would have if these keys were “thrown” into  $A$  uniformly at random.

Rezultat za stranicu 435

## 10.2. Hash Tables 413

collisions for common groups of strings. In particular, “temp01” and “temp10” collide using this function, as do “stop”, “tops”, “pots”, and “spot”. A better hash code should somehow take into consideration the positions of the  $x_i$ ’s. An alternative hash code, which does exactly this, is to choose a nonzero constant,  $a$ , and use as a hash code the value

$$x_0a^{n-1} + x_1a^{n-2} + \cdots + x_{n-2}a + x_{n-1}.$$

Mathematically speaking, this is simply a polynomial in  $a$  that takes the components  $(x_0, x_1, \dots, x_{n-1})$  of an object  $x$  as its coefficients. This hash code is therefore called a polynomial hash code. By Horner’s rule (see Exercise C-3.50), this polynomial can be computed as

$$x_{n-1} + a(x_{n-2} + a(x_{n-3} + \cdots + a(x_2 + a(x_1 + ax_0)) \cdots)).$$

Intuitively, a polynomial hash code uses multiplication by different powers as a way to spread out the influence of each component across the resulting hash code.

Of course, on a typical computer, evaluating a polynomial will be done using the finite bit representation for a hash code; hence, the value will periodically overflow the bits used for an integer. Since we are more interested in a good spread of the object  $x$  with respect to other keys, we simply ignore such overflows. Still, we

should be mindful that such overflows are occurring and choose the constant so that it has some nonzero, low-order bits, which will serve to preserve some of the information content even as we are in an overflow situation.

We have done some experimental studies that suggest that 33, 37, 39, and 41 are particularly good choices for  $a$  when working with character strings that are English words. In fact, in a list of over 50,000 English words formed as the union of the word lists provided in two variants of Unix, we found that taking  $a$  to be 33, 37, 39, or 41 produced less than 7 collisions in each case!

### Cyclic-Shift **Hash** Codes

A variant of the polynomial **hash** code replaces multiplication by  $a$  with a cyclic shift of a partial sum by a certain number of bits. For example, a 5-bit cyclic shift of the 32-bit value 00111

101100101101010100010101000 is achieved by taking

the leftmost five bits and placing those on the rightmost side of the representation, resulting in 10110010110101010001010100000111

. While this operation has little

natural meaning in terms of arithmetic, it accomplishes the goal of varying the bits of the calculation. In Python, a cyclic shift of bits can be accomplished through careful use of the bitwise operators  $\ll$  and  $\gg$ , taking care to truncate results to 32-bit integers.

Rezultat za stranicu 443

## 10.2. **Hash** Tables 421

### Operation

#### List

#### **Hash** Table



expected

worst case

getitem

$O(n)$

$O(1)$

$O(n)$

setitem

$O(n)$

$O(1)$

$O(n)$

delitem

$O(n)$

$O(1)$

$O(n)$

len

$O(1)$

$O(1)$

$O(1)$

iter

$O(n)$

$O(n)$

$O(n)$

Table 10.2: Comparison of the running times of the methods of a map realized by

means of an unsorted list (as in Section 10.1.5) or a **hash** table. We let  $n$  denote

the number of items in the map, and we assume that the bucket array supporting

the **hash** table is maintained such that its capacity is proportional to the number of items in the

map.

In practice, **hash** tables are among the most efficient means for implementing a map, and it is essentially taken for granted by programmers that their core operations run in constant time. Python's dictclass is implemented with hashing, and the Python interpreter relies on dictionaries to retrieve an object that is referenced by an identifier in a given namespace. (See Sections 1.10 and 2.5.) The basic command `c=a+b` involves two calls to

`getitem`

in the dictionary for the local

namespace to retrieve the values identified as `a` and `b`, and a call to

`setitem`

to store the result associated with name `c` in that namespace. In our own algorithm

analysis, we simply presume that such dictionary operations run in constant time, independent of the number of entries in the namespace. (Admittedly, the number of entries in a typical namespace can almost surely be bounded by a constant.)

In a 2003 academic paper [31], researchers discuss the possibility of exploiting

a **hash** table's worst-case performance to cause a denial-of-service (DoS) attack of Internet technologies. For many published algorithms that compute **hash** codes,

they note that an attacker could precompute a very large number of moderate-length

strings that all **hash** to the identical 32-bit **hash** code. (Recall that by any of the hashing schemes we describe, other than double hashing, if two keys are mapped to the same **hash** code, they will be inseparable in the collision resolution.)

In late 2011, another team of researchers demonstrated an implementation of

just such an attack [61]. Web servers allow a series of key-value parameters to be embedded in a URL using a syntax such as `?key1=val1&key2=val2&key3=val3`.

Typically, those key-value pairs are immediately stored in a map by the server,

and a limit is placed on the length and number of such parameters presuming that storage time in the map will be linear in the number of entries. If all keys were to collide, that storage requires quadratic time (causing the server to perform an inordinate amount of work). In spring of 2012, Python developers distributed a security patch that introduces randomization into the computation of `hash` codes for strings, making it less tractable to reverse engineer a set of colliding strings.

Rezultat za stranicu 442

## 420 Chapter 10. Maps, `Hash` Tables, and Skip Lists

### 10.2.3 Load Factors, Rehashing, and Efficiency

In the `hash` table schemes described thus far, it is important that the load factor,  $\lambda = n/N$ , be kept below 1. With separate chaining, as  $\lambda$  gets very close to 1, the probability of a collision greatly increases, which adds overhead to our operations, since we must revert to linear-time list-based methods in buckets that have collisions. Experiments and average-case analyses suggest that we should maintain  $\lambda < 0.9$  for `hash` tables with separate chaining.

With open addressing, on the other hand, as the load factor  $\lambda$  grows beyond 0.5 and starts approaching 1, clusters of entries in the bucket array start to grow as well. These clusters cause the probing strategies to “bounce around” the bucket array for a considerable amount of time before they find an empty slot. In Exercise C-10.36, we explore the degradation of quadratic probing when  $\lambda \geq 0.5$ . Experiments suggest that we should maintain  $\lambda < 0.5$  for an open addressing scheme with linear probing, and perhaps only a bit higher for other open addressing schemes (for example, Python’s implementation of open addressing enforces that  $\lambda < 2/3$ ).

If an insertion causes the load factor of a `hash` table to go above the specified threshold, then it is common to resize the table (to regain the specified load factor)

and to reinsert all objects into this new table. Although we need not define a new `hash` code for each object, we do need to reapply a new compression function that takes into consideration the size of the new table. Each rehashing will generally scatter the items throughout the new bucket array. When rehashing to a new table, it is a good requirement for the new array's size to be at least double the previous size.

Indeed, if we always double the size of the table with each rehashing operation, then we can amortize the cost of rehashing all the entries in the table against the time used to insert them in the first place (as with dynamic arrays; see Section 5.3).

### Efficiency of `Hash` Tables

Although the details of the average-case analysis of hashing are beyond the scope of this book, its probabilistic basis is quite intuitive. If our `hash` function is good, then we expect the entries to be uniformly distributed in the  $N$  cells of the bucket

array. Thus, to store  $n$  entries, the expected number of keys in a bucket would be  $\lceil n/N \rceil$ , which is  $O(1)$  if  $n$  is  $O(N)$ .

The costs associated with a periodic rehashing, to resize a table after occasional insertions or deletions can be accounted for separately, leading to an additional  $O(1)$  amortized cost for

`setitem`

and

`getitem`

.

In the worst case, a poor `hash` function could map every item to the same bucket.

This would result in linear-time performance for the core map operations with separate chaining, or with any open addressing model in which the secondary sequence of probes depends only on the `hash` code. A summary of these costs is given in

Table 10.2.

## 452 Chapter 10. Maps, **Hash** Tables, and Skip Lists

### 10.6 Exercises

For help with exercises, please visit the site, [www.wiley.com/college/goodrich](http://www.wiley.com/college/goodrich).

#### Reinforcement

R-10.1 Give a concrete implementation of the `pop` method in the context of the `MutableMapping` class, relying only on the five primary abstract methods of that class.

R-10.2 Give a concrete implementation of the `items()` method in the context of the `MutableMapping` class, relying only on the five primary abstract methods of that class. What would its running time be if directly applied to the `UnsortedTableMap` subclass?

R-10.3 Give a concrete implementation of the `items()` method directly within the `UnsortedTableMap` class, ensuring that the entire iteration runs in  $O(n)$  time.

R-10.4 What is the worst-case running time for inserting  $n$  key-value pairs into an initially empty map  $M$  that is implemented with the `UnsortedTableMap` class?

R-10.5 Reimplement the `UnsortedTableMap` class from Section 10.1.5, using the `PositionalList` class from Section 7.4 rather than a Python list.

R-10.6 Which of the **hash** table collision-handling schemes could tolerate a load factor above 1 and which could not?

R-10.7 Our `Position` classes for lists and trees support the `eq` method so that

two distinct position instances are considered equivalent if they refer to the same underlying node in a

structure. For positions to be allowed as keys in a **hash** table, there must be a definition for the

**hash**

method that

is consistent with this notion of equivalence. Provide such a

**hash**

method.

R-10.8 What would be a good **hash** code for a vehicle identification number that is a string of numbers and letters of the form “9X9XX99X9XX999999,” where a “9” represents a digit and an “X” represents a letter?

R-10.9 Draw the 11-entry **hash** table that results from using the **hash** function,  $h(i) = (3i + 5) \bmod 11$ , to **hash** the keys 12, 44, 13, 88, 23, 94, 11, 39, 20, 16, and 5, assuming collisions are handled by chaining.

R-10.10 What is the result of the previous exercise, assuming collisions are handled by linear probing?

R-10.11 Show the result of Exercise R-10.9, assuming collisions are handled by quadratic probing, up to the point where the method fails.

Rezultat za stranicu 432

## 410 Chapter 10. Maps, **Hash** Tables, and Skip Lists

### 10.2 **Hash** Tables

In this section, we introduce one of the most practical data structures for implementing a map, and the one that is used by Python’s own implementation of the `dict` class. This structure is known as a **hash** table.

Intuitively, a map  $M$  supports the abstraction of using keys as indices with a syntax such as  $M[k]$ . As a mental warm-up, consider a restricted setting in which a map with  $n$  items uses keys that are known to be integers in a range from 0 to  $N-1$  for some  $N \geq n$ . In this case, we can represent the map using a lookup table

of length N, as diagrammed in Figure 10.3.

0 123456789 1 0  
DZ C Q

Figure 10.3: A lookup table with length 11 for a map containing items (1,D), (3,Z), (6,C), and (7,Q).

In this representation, we store the value associated with key *kat* at index *kof* of the table (presuming that we have a distinct way to represent an empty slot). Basic map operations of `getitem`

,  
`setitem`

, and

`delitem`

can be implemented in

$O(1)$  worst-case time.

There are two challenges in extending this framework to the more general setting of a map. First, we may not wish to devote an array of length  $N$  if it is the case that  $N \gg \text{number of keys}$ . Second, we do not in general require that a map's keys be integers.

The novel concept for a `hash` table is the use of a `hash` function to map general keys to corresponding indices in a table. Ideally, keys will be well distributed in the range from 0 to  $N-1$  by a `hash` function, but in practice there may be two or more distinct keys that get mapped to the same index. As a result, we will conceptualize our table as a bucket array, as shown in Figure 10.4, in which each bucket may manage a collection of items that are sent to a specific index by the `hash` function. (To save space, an empty bucket may be replaced by `None`.)

0 123456789 1 0  
(1,D) (25,C)

(3,F)

(14,Z)(39,C)(6,A) (7,Q)

Figure 10.4: A bucket array of capacity 11 with items (1,D), (25,C), (3,F), (14,Z), (6,A), (39,C), and (7,Q), using a simple **hash** function.

Rezultat za stranicu 439

## 10.2. **Hash** Tables 417

### 10.2.2 Collision-Handling Schemes

The main idea of a **hash** table is to take a bucket array,  $A$ , and a **hash** function,  $h$ , and

use them to implement a map by storing each item  $(k,v)$  in the “bucket”  $A[h(k)]$ .

This simple idea is challenged, however, when we have two distinct keys,  $k_1$  and  $k_2$ , such that  $h(k_1)=h(k_2)$ . The existence of such collisions prevents us from simply inserting a new item  $(k,v)$  directly into the bucket  $A[h(k)]$ . It also complicates our procedure for performing insertion, search, and deletion operations.

#### Separate Chaining

A simple and efficient way for dealing with collisions is to have each bucket  $A[j]$  store its own secondary container, holding items  $(k,v)$  such that  $h(k)=j$ . A natural choice for the secondary container is a small map instance implemented using a list, as described in Section 10.1.5. This collision resolution rule is known as separate chaining, and is illustrated in Figure 10.6.

A123456789 1 0 01 112

123825

9054

28413618 10

Figure 10.6: A **hash** table of size 13, storing 10 items with integer keys, with colli-



sions resolved by separate chaining. The compression function is  $h(k) = k \bmod 13$ .

For simplicity, we do not show the values associated with the keys.

In the worst case, operations on an individual bucket take time proportional to the size of the bucket. Assuming we use a good `hash` function to index the items of our map in a bucket array of capacity  $N$ , the expected size of a bucket is  $n/N$ .

Therefore, if given a good `hash` function, the core map operations run in  $O(\lceil n/N \rceil)$ .

The ratio  $\lambda = n/N$ , called the load factor of the `hash` table, should be bounded by a small constant, preferably below 1. As long as  $\lambda = O(1)$ , the core operations on the `hash` table run in  $O(1)$  expected time.