

2 Chapter 1. Python Primer

1.1 Python Overview

Building data structures and algorithms requires that we communicate detailed instructions to a computer. An excellent way to perform such communications is using a high-level computer language. A language was originally developed by Guido van Rossum in the early 1990s, and has since become a prominently used language in industry and education. The second major version of the language (more specifically, Python 3.1 or later). The latest version of the language is freely available at www.python.org, along with documentation and tutorials.

In this chapter, we provide an overview of the Python programming language, and we continue this discussion in the next chapter, focusing on object-oriented principles. We assume that readers are familiar with programming in some language, although not necessarily using Python. This book does not provide a complete description of the Python language (there are numerous language references for that purpose), but it does provide a detailed overview.

1.1.1 The Python Interpreter

Python is formally an interpreted language. Commands are executed through a piece of software known as the Python interpreter. The interpreter receives a command, evaluates that command, and reports the result of the command. While the interpreter can be used interactively, it can also be used to execute scripts. Scripts are files that are typically stored in a file named with the .py suffix (e.g., demo.py).

On most operating systems, the Python interpreter can be started by typing `python` from the command line. By default, the interpreter starts in interactive mode with a clean workspace. Commands from a predefined script saved in a file (e.g., demo.py) are executed by providing an argument (e.g., `python demo.py`), or using an additional `-i` flag to enter interactive mode (e.g., `python -i demo.py`).

Many integrated development environments (IDEs) provide richer software development platforms for Python, including one named IDLE that is included with the standard Python distribution. IDLE provides a graphical user interface and step-by-step execution of a program while examining key variable values.

1.2. Objects in Python 11

The set and frozenset Classes

Python's set class represents the mathematical notion of a set, namely a collection of elements, without duplicates, and without an inherent order to those elements.

The major advantage of using a set, as opposed to a list, is that it has a highly

optimized method for checking whether a specific element is contained in the set. This is based on a data structure

(topic of Chapter 10). However, there are two important restrictions due to the

algorithmic underpinnings. The first is that the set does not maintain the elements

in any particular order. The second is that only instances of immutable types can be

added to a Python set. Therefore, objects such as integers, floating-point numbers,

and character strings are eligible to be elements of a set. It is possible to maintain a

set of tuples, but not a set of lists or a set of sets, as lists and sets are mutable. The

frozenset class is an immutable form of the set type, so it is legal to have a set of

frozensets.

Python uses curly braces {} as delimiters for a set, for example, as {1, 7}

or {

red

,

green

,

blue

}. The exception to this rule is that {} does not

represent an empty set; for historical reasons, it represents an empty dictionary (see next paragraph). Instead, the

set constructor set() takes an iterable parameter. If an iterable parameter is sent to the constructor, then the set of distinct elements

is produced. For example, set(

hello

) produces {

h

,

e

,

l

,

o

}. The

The dict Class

Python's dict class represents a dictionary, a mapping from a set of distinct keys

to associated values. For example, a dictionary might map from unique student ID

numbers, to larger student records (such as the student's name, address, and course

grades). Python implements a dict using an almost identical approach to that of a

set, but with storage of the associated values.

A dictionary literal also uses curly braces, and because dictionaries were intro-

duced in Python prior to sets, the literal form {} produces an empty dictionary.

A nonempty dictionary is expressed using a comma-separated series of key:value pairs. For example, the dictio-

ga

:

Irish

,

de

:

644 Chapter 14. Graph Algorithms

14.3.2 DFS Implementation and Extensions

We begin by providing a **Python** implementation of the basic depth-first search algorithm, originally described with pseudo-code in Code Fragment 14.4. Our DFS function is presented in Code Fragment 14.5.

```
1 def DFS(g, u, discovered):
2     ... Perform DFS of the undiscovered portion of Graph g starting at Vertex u.
3
4     discovered is a dictionary mapping each vertex to the edge that was used to
5     discover it during the DFS. (u should be 'discovered' prior to the call.)
6     Newly discovered vertices will be added to the dictionary as a result.
7
8     for e in g.incident
9         edges(u): # for every outgoing edge from u
10            v = e.opposite(u)
11            if v not in discovered: # v is an unvisited vertex
12                discovered[v] = e # e is the tree edge that discovered v
13            DFS(g, v, discovered) # recursively explore from v
```

Code Fragment 14.5: Recursive implementation of depth-first search on a graph, starting at a designated vertex u.

In order to track which vertices have been visited, and to build a representation of the resulting DFS tree, our implementation introduces a third parameter, named `discovered`. This parameter serves as a **dictionary** mapping each vertex to the tree edge that was used to discover that vertex. As a technicality, we assume that the source vertex `u` is already in the `discovered` dictionary. Thus, a caller might start the traversal as follows:

```
result = {u: None} # a new dictionary, with u trivially discovered
DFS(g, u, result)
```

The **dictionary** serves two purposes. Internally, the **dictionary** provides a mechanism for recognizing visited vertices. At the conclusion of the process, the values within the **dictionary** are the DFS tree edges.

Because the **dictionary** is hash-based, the test, `if v not in discovered`, and the record-keeping step, `discovered[v] = e`, run in $O(1)$ expected time, rather than worst-case time. In practice, this is a compromise we are willing to accept, but it does violate the formal analysis. A better approach could be used as indices into an array-based lookup table rather than a hash-based map. Alternatively, we could store the tree edge directly as part of the vertex instance.

14.2.5 Python Implementation

In this section, we provide an implementation of the Graph ADT. Our implementa-

tion will support directed or undirected graphs, but for ease of explanation, we describe it in the context of an

We use a variant of the adjacency map representation. For each vertex v , we

use a Python dictionary to represent the secondary incidence map $I(v)$. However,

we do not explicitly maintain lists V and E , as originally described in the edge list

representation. The list V is replaced by a top-level dictionary D that maps each

vertex v to its incidence map $I(v)$; note that we can iterate through all vertices by

generating the set of keys for dictionary D . By using such a dictionary D to map

vertices to the secondary incidence maps, we need not maintain references to those incidence maps as part of the

$O(1)$ expected time. This greatly simplifies our implementation. However, a

consequence of our design is that some of the worst-case running time bounds for the graph ADT operations, given

than maintain list E , we are content with taking the union of the edges found in the

various incidence maps; technically, this runs in $O(n+m)$ time rather than strictly

$O(m)$ time, as the dictionary D has n keys, even if some incidence maps are empty.

Our implementation of the graph ADT is given in Code Fragments 14.1 through

14.3. Classes `Vertex` and `Edge`, given in Code Fragment 14.1, are rather simple,

and can be nested within the more complex `Graph` class. Note that we define the

hash

method for both `Vertex` and `Edge` so that those instances can be used as

keys in Python's hash-based sets and dictionaries. The rest of the `Graph` class is

given in Code Fragments 14.2 and 14.3. Graphs are undirected by default, but can be declared as directed with a

Internally, we manage the directed case by having two different top-level dictionary

instances,

outgoing and

incoming, such that

outgoing[v] maps to another

dictionary representing $\text{out}(v)$, and

incoming[v] maps to a representation of $\text{in}(v)$.

In order to unify our treatment of directed and undirected graphs, we continue to use the

outgoing and

incoming identifiers in the undirected case, yet as aliases

to the same dictionary. For convenience, we define a utility named `is`

directed to

allow us to distinguish between the two cases.

For methods `degree` and `incident`

edges, which each accept an optional param-

eter to differentiate between the outgoing and incoming orientations, we choose the appropriate map before processing

vertex v , we always initial-

ize

outgoing[v] to an empty dictionary for new vertex v . In the directed case, we

independently initialize

incoming[v] as well. For the undirected case, that step is

unnecessary as

outgoing and

incoming are aliases. We leave the implementations

of methods `remove`

vertex and `remove`

edge as exercises (C-14.37 and C-14.38).

40 Chapter 1. Python Primer

We see lazy evaluation used in many of Python's libraries. For example, the

`dictionary` class supports methods `keys()`, `values()`, and `items()`, which respectively produce a view of all keys, values, or (key,value) pairs within a `dictionary`.

None of these methods produces an explicit list of results. Instead, the views that are produced are iterable objects. Calling the list class constructor with the iteration as a parameter. For example, the syntax `list(range(1000))` produces a list instance with values from 0 to 999, while the syntax `list(d.values())` produces a list that has elements based upon the current values of `dictionary` `d`. We can similarly construct a tuple or set instance based upon a given iterable.

Generators

In Section 2.3.4, we will explain how to define a class whose instances serve as iterators. However, the most convenient technique for creating iterators in Python

is through the use of generators. A generator is implemented with a syntax that is very similar to a function, but instead of returning values, a `yield` statement is

executed to indicate each element of the series. As an example, consider the goal

of determining all factors of a positive integer. For example, the number 100 has factors 1, 2, 4, 5, 10, 20, 25, 50, return a list containing all factors, implemented as:

```
def factors(n): # traditional function that computes factors
```

```
    results = [] # store factors in a new list
```

```
    for k in range(1, n+1):
```

```
        if n % k == 0: # divides evenly, thus k is a factor
```

```
            results.append(k) # add k to the list of factors
```

```
    return results # return the entire list
```

In contrast, an implementation of a generator for computing those factors could be implemented as follows:

```
def factors(n): # generator that computes factors
```

```
    for k in range(1, n+1):
```

```
        if n % k == 0: # divides evenly, thus k is a factor
```

```
            yield k # yield this factor as next result
```

Notice use of the keyword `yield` rather than `return` to indicate a result. This indicates to Python that we are defining a generator, rather than a traditional function. It is illegal to combine `yield` and `return` statements in the same implementation, other than a zero-argument `return` statement to cause a generator to end its execution. If a programmer writes a loop such as `for factor in factors(100):`, an instance of our generator is created. For each iteration of the loop, Python executes our procedure

14 Chapter 1. Python Primer

Python carefully extends the semantics of // and % to cases where one or both operands are negative. For the sake of notation, let us assume that variables n and m represent respectively the dividend and divisor of a quotient n/m .

Let q and r be integers such that

$q = n // m$ and $r = n \% m$. Python guarantees that q

$m + r$ will equal n . We

already saw an example of this identity with positive operands, as $6 \cdot 4 + 3 = 27$.

When the divisor m is positive, Python further guarantees that $0 \leq r < m$. As

a consequence, we find that $27 // 4$ evaluates to 6 and $27 \% 4$ evaluates

to 3 , as $(6 \cdot 4) + 3 = 27$. When the divisor is negative, Python guarantees that

$m < r \leq 0$. As an example, $27 // -4$ is -7 and $27 \% -4$ is 1 , satisfying the

identity $27 = (-7) \cdot (-4) + (1)$.

The conventions for the // and % operators are even extended to floating-

point operands, with the expression $q = n // m$ being the integral floor of the

quotient, and $r = n \% m$ being the remainder, to ensure that q

$m + r$ equals

n . For example, $8.2 // 3.14$ evaluates to 2.0 and $8.2 \% 3.14$ evaluates to 1.92 , as

$2.0 \cdot 3.14 + 1.92 = 8.2$.

Bitwise Operators

Python provides the following bitwise operators for integers:

- bitwise complement (pre- x unary operator)

- & bitwise and

- | bitwise or

- bitwise exclusive-or

- << shift bits left, filling in with zeros

- >> shift bits right, filling in with sign bit

Sequence Operators

Each of Python's built-in sequence types (str, tuple, and list) support the following operator syntaxes:

$s[j]$ element at index j

$s[start:stop]$ slice including indices $[start, stop)$

$s[start:stop:step]$ slice including indices $start, start + step, start + 2$

$step, \dots$, up to but not including or stop

$s+t$ concatenation of sequences

k

s shorthand for $s+s+s+\dots$ (k times)

val in s containment check

val not in s non-containment check

Python relies on zero-indexing of sequences, thus a sequence of length n has elements indexed from 0 to $n-1$ inclusive. Python also supports the use of negative indices, which denote a distance from the end of the sequence; index -1 denotes the last element, index -2 the second to last, and so on. Python uses a slicing

15.1. Memory Management 703

15.1.3 Additional Memory Used by the Python Interpreter

We have discussed, in Section 15.1.1, how the Python interpreter allocates memory for objects within a memory heap. However, this is not the only memory that is used when executing a Python program.

The Run-Time Call Stack

Stacks have a most important application to the run-time environment of Python programs. A running Python program has a private stack, known as the call stack or Python interpreter stack, that is used to keep track of the nested sequence of currently active (that is, nonterminated) invocations of functions. Each entry of the stack is a structure known as an activation record or frame, storing important information about an invocation of a function.

At the top of the call stack is the activation record of the running call, that is, the function activation that currently has control of the execution. The remaining elements of the stack are activations that have invoked another function and are currently waiting for that other function to return control when it terminates. The order of the elements in the stack corresponds to the chain of invocations of the currently active functions. When a new function is called, an activation record is created and added to the stack. When the function terminates, the Python interpreter resumes the processing of the previously suspended call.

Each activation record includes a dictionary representing the local namespace for the function call. (See Sections 1.10 and 2.5 for further discussion of namespaces). The namespace maps identifiers, which serve as parameters and local variables, to object values, although the objects being referenced still reside in the memory heap. The activation record for a function call also includes a reference to the function definition itself, and a reference to maintain the address of the statement within the function that is currently executing. When one function returns control to another, the stored program counter for the suspended function allows the interpreter to properly continue execution of that function.

Implementing Recursion

One of the benefits of using a stack to implement the nesting of function calls is that it allows programs to use recursion, as discussed in Chapter 4. We implicitly described the concept of the call stack and the use of activation records within our portrayal of recursion traces in

1.1. Python Overview 3

1.1.2 Preview of a Python Program

As a simple introduction, Code Fragment 1.1 presents a Python program that computes the grade-point average (GPA) for a student based on letter grades that are entered by a user. Many of the Python's syntax relies heavily on the use of whitespace. Individual statements are typically concluded with a newline character, although a command can extend to another line, either with a colon character if the command has not yet been closed, such as the `map` command in defining value `map`.

Whitespace is also key in delimiting the bodies of control structures in Python. Specifically, a block of code is indented to designate it as the body of a control structure, and nested control structures, including a nested conditional structure.

Comments are annotations provided for human readers, yet ignored by the Python interpreter. The primary syntax for comments in Python is based on use of the `#` character, which designates a comment line.

```
print(
    Welcome to the GPA calculator.
)
print(
    Please enter all your letter grades, one per line.
)
print(
    Enter a blank line to designate the end.
)
# map from letter grade to point value
points = {
    A+
:4.0,
    A
:4.0,
    A-
:3.67,
    B+
:3.33,
    B
:3.0,
    B-
:2.67,
    C+
:2.33,
    C
:2.0,
    C-
:1.67,
    D+
:1.33,
    D
:1.0,
    F
:0.0}
num
courses = 0
```


46 Chapter 1. Python Primer

1.10 Scopes and Namespaces

When computing a sum with the syntax `x+y` in Python, the names `x` and `y` must have been previously associated with objects that serve as values; a `NameError` will be raised if no such definitions are found. The process of determining the value associated with an identifier is known as name resolution.

Whenever an identifier is assigned to a value, that definition is made with a specific scope. Top-level assignments are typically made in what is known as global scope. Assignments made within the body of a function typically have scope that is local to that function call. Therefore, an assignment, `x=5`, within a function has no effect on the identifier, `x`, in the broader scope.

Each distinct scope in Python is represented using an abstraction known as a namespace. A namespace manages all identifiers that are currently defined in a given scope. Figure 1.8 portrays two namespaces, one being that of a caller to our `count` function from Section 1.5, and the other being the local namespace during the execution of that function.

The figure shows two namespaces side-by-side. The left namespace (caller's) contains: 'str', 'A', 'str', 'CS', '-oat', '3.56int', '2', 'itemdatagrades', 'majorgp', 'targetn', 'list', 'str', 'B+', 'str', 'A-', 'str'. The right namespace (local scope) contains: 'A', 'str', 'B+', 'str', 'A-', 'str'.

Figure 1.8: A portrayal of the two namespaces associated with a user's call `count(grades,`

`A`), as defined in Section 1.5. The left namespace is the caller's and the right namespace represents the local scope of the function.

Python implements a namespace with its own dictionary that maps each identifying string (e.g.,

`n`) to its associated value. Python provides several ways to examine a given namespace. The function, `dir`, reports the names of the identifiers in a given namespace (i.e., the keys of the dictionary), while the function, `vars`, returns the full dictionary. By default, calls to `dir()` and `vars()` report on the most locally enclosing namespace in which they are executed.

12 Chapter 1. Python Primer

1.3 Expressions, Operators, and Precedence

In the previous section, we demonstrated how names can be used to identify existing objects, and how literals and constructors can be used to create instances of built-in classes. Existing values can be combined into larger syntactic expressions using a variety of special symbols and keywords known as operators. The semantics of an operator depends upon the type of its operands. For example, when `a` and `b` are numbers, the syntax `a+b` indicates addition, while if `a` and `b` are strings, the operator indicates concatenation. In this section, we describe Python's operators in various contexts of the language. We continue, in Section 1.3.1, by discussing compound expressions, such as

`a+b`

`c`, which rely on the evaluation of two or more operations. The order

in which the operations of a compound expression are evaluated can affect the overall value of the expression. For example, the precedence for evaluating operators, and it allows a programmer to override this order by using explicit parentheses to group subexpressions.

Logical Operators

Python supports the following keyword operators for Boolean values:

`not` unary negation

`and` conditional and

`or` conditional or

The `and` and `or` operators short-circuit, in that they do not evaluate the second operand if the result can be determined based on the value of the first operand.

This feature is useful when constructing Boolean expressions in which we first test that a certain condition holds (such as a reference not being `None`), and then test a condition that could have otherwise generated an error condition had the prior test not succeeded.

Equality Operators

Python supports the following operators to test two notions of equality:

`is` same identity

`is not` different identity

`==` equivalent

`!=` not equivalent

The expression `a is b` evaluates to `True`, precisely when identifiers `a` and `b` are aliases for the same object. The expression `a == b` tests a more general notion of equivalence. If identifiers `a` and `b` refer to the same object, then `a == b` should also evaluate to `True`. Yet `a == b` also evaluates to `True` when the identifiers refer to