

10.4. Skip Lists 437

10.4 Skip Lists

An interesting data structure for realizing the sorted map ADT is the **skip list**. In Section 10.3.1, we saw that a sorted array will allow $O(\log n)$ -time searches via the binary search algorithm. Unfortunately, update operations on a sorted array have $O(n)$ worst-case running time because of the need to shift elements. In Chapter 7 we demonstrated that linked lists support very efficient update operations, as long as the position within the list is identified. Unfortunately, we cannot perform fast searches on a standard linked list; an efficient means for direct accessing an element of a sequence by index.

Skip lists provide a clever compromise to efficiently support search and update operations. A **skip list** S for a map M consists of a series of lists $\{S_0, S_1, \dots, S_h\}$.

Each list S_i stores a subset of the items of M sorted by increasing keys, plus items with two sentinel keys denoted $-\infty$ and $+\infty$, where $-\infty$ is smaller than every possible key that can be inserted in M and $+\infty$ is larger than every possible key that can be inserted in M . In addition, the lists in S satisfy the following:

- List S_0 contains every item of the map M (plus sentinels $-\infty$ and $+\infty$).
- For $i=1, \dots, h-1$, list S_i contains (in addition to $-\infty$ and $+\infty$) a randomly generated subset of the items in list S_{i-1} .
- List S_h contains only $-\infty$ and $+\infty$.

An example of a **skip list** is shown in Figure 10.10. It is customary to visualize a **skip list** S with list S_0 at the bottom and lists S_1, \dots, S_h above it. Also, we refer to h as the height of **skip list** S .

Intuitively, the lists are set up so that S_{i+1} contains more or less alternate items of S_i . As we shall see in the details of the insertion method, the items in S_{i+1} are chosen at random from the items in S_i by picking each item from S_i to also be in S_{i+1} with probability $1/2$. That is, in essence, we flip a coin for each item in S_i .

3125

25

--

--1717

17

17 12 S5

S4

S3

S2

S1

S0 55

55

55

55 12 17 20 25 31 38 39 44 50 +--+--+--+--

44 38 31 25

Figure 10.10: Example of a **skip list** storing 10 items. For simplicity, we show only the items' keys, not their associated values.

438 Chapter 10. Maps, Hash Tables, and Skip Lists

and place that item in S_{i+1} if the coin comes up ·heads·. Thus, we expect S_1 to have about $n/2$ items, S_2 to have about $n/4$ items, and, in general, S_i to have about $n/2^i$ items. In other words, we expect the height of S to be about $\log n$. The halving of the number of items from one list to the next is not enforced as an explicit property of skip lists, however. Instead, randomization is used.

Functions that generate numbers that can be viewed as random numbers are built into most modern computers, because they are used extensively in computer games, cryptography, and computer simulations. Random number generators, generate random-like numbers, starting with an initial seed. (See discussion of random module in Section 1.11.1.) Other methods use hardware devices to extract ·true· random numbers from nature. In any case, we will assume that our computer has a good random number generator. The main advantage of using randomization in data structure and algorithm design is that the structures and functions that result are usually simple and efficient.

The skip list has the same logarithmic time bounds for searching as is achieved by the binary search algorithm, yet it extends that performance to update methods when inserting or deleting items. skip list, while binary search has a worst-case bound with a sorted table.

A skip list makes random choices in arranging its structure in such a way that search and update times are $O(\log n)$ on average, where n is the number of items in the map. Interestingly, the notion of average time complexity used here does not depend on the probability distribution used to help decide where to place the new item. The running time is averaged over all possible outcomes of the random numbers used when inserting entries.

Using the position abstraction used for lists and trees, we view a skip list as a two-dimensional collection of positions arranged horizontally into levels and vertically into towers. Each level is a list S_i

and each tower contains positions storing

the same item across consecutive lists. The positions in a skip list can be traversed using the following operations:

$\text{next}(p)$: Return the position following p on the same level.

$\text{prev}(p)$: Return the position preceding p on the same level.

$\text{below}(p)$: Return the position below p in the same tower.

$\text{above}(p)$: Return the position above p in the same tower.

We conventionally assume that the above operations return `None` if the position

requested does not exist. Without going into the details, we note that we can easily implement a skip list by means of a doubly linked list structure. The structure is essentially a collection of doubly linked lists aligned at towers, which are also doubly linked lists.

Removal in a Skip List

Like the search and insertion algorithms, the removal algorithm for a skip list is quite simple. In fact, it is even easier than the insertion algorithm. That is, to perform the map operation `del M[k]` we begin by executing method `SkipSearch(k)`. If the position `p` stores an entry with key different from `k`, we raise a `KeyError`.

Otherwise, we remove `p` and all the positions above `p`, which are easily accessed by using above operations to climb up the tower of this entry in `S` starting at position `p`. While removing levels of the tower, we reestablish links between the horizontal neighbors of each removed position. The removal algorithm is illustrated in Figure 10.13 and a detailed de-

$O(\log n)$ expected running time. Before we give this analysis, however, there are some minor improvements to the skip-list data structure we would like to discuss. First, we do not actually need to store references to values at more efficiently represent a tower as a single object, storing the key-value pair, and maintaining `jprevious` references and `jnext` references if the tower reaches level `S`.

j. Second, for the horizontal axes, it is possible to keep the list singly linked, storing only the next references. We can perform insertions and removals in strictly a top-down, scan-forward fashion. Exercise C-10.44. Neither of these optimizations improve the asymptotic performance of skip lists by more than a constant factor, but these improvements can, nevertheless, be meaningful in practical applications, which are discussed in Chapter 11.

```

31 S5
S4
S3
S2
S1---
-- 1212 --
1717 25
25 31
3142
55 5055+-
+-+
+-
+-----
17
38
38 39 424242
44
445555+-
17
17
20 2525
S0

```

Figure 10.13: Removal of the entry with key 25 from the skip list of Figure 10.12. The positions visited after the search for the position of `S0` holding the entry are highlighted. The positions removed are drawn with dashed lines.

Algorithm SkipSearch(k) :

Input: A search key k

Output: Position p in the bottom list S₀ with the largest key such that key(p) ≤ k

p ← start {begin at start position }

while below(p) ≠ None do

p ← below(p) {drop down }

while k < key(next(p)) do

p ← next(p) {scan forward }

return p.

Code Fragment 10.12: Algorithm to search a skip list S for key k.

As it turns out, the expected running time of algorithm SkipSearch on a skip list with n entries is $O(\log n)$. We postpone the justification of this fact, however, until after we discuss the implementation of the update methods for skip lists. Navigation starting at the position identified by SkipSearch(k) can be easily used to provide the additional forms of searches in the sorted map ADT (e.g., find range).

Insertion into a Skip List

The execution of the map operation $M[k] = v$ begins with a call to SkipSearch(k).

This gives us the position p of the bottom-level item with the largest key less than or equal to k (note that p may hold the special item with key ∞). If key(p) = k, the

associated value is overwritten with v. Otherwise, we need to create a new tower for item (k, v). We insert (k, v) immediately after position p within S₀. After inserting

the new item at the bottom level, we use randomization to decide the height of the tower for the new item. We flip a coin, and if the coin comes up tails, then we stop here. Else (the coin comes up heads), we backtrack to the previous (next higher)

level and insert (k, v) in this level at the appropriate position. We again flip a coin;

if it comes up heads, we go to the next higher level and repeat. Thus, we continue to insert the new item (k, v) in list

We link together all the references to the new item (k, v) created in this process to

create its tower. A coin flip can be simulated with Python's built-in pseudo-random number generator from the random module, which returns 0 or 1, each with probability 1/2.

We give the insertion algorithm for a skip list S in Code Fragment 10.13 and

we illustrate it in Figure 10.12. The algorithm uses an insertAfterAbove(p, q, (k, v))

method that inserts a position storing the item (k, v) after position p (on the same

level as p) and above position q, returning the new position r (and setting internal references so that next, prev, above, and below methods will work correctly for p,

q, and r). The expected running time of the insertion algorithm on a skip list with n entries is $O(\log n)$, which we show in Section 10.4.2.

444 Chapter 10. Maps, Hash Tables, and Skip Lists

Bounding the Height of a Skip List

Because the insertion step involves randomization, a more accurate analysis of skip lists involves a bit of probability. At first, this might seem like a major undertaking, for a complete and thorough probabilistic analysis could require deep mathematics (and, indeed, there are several ways to understand the expected asymptotic behavior of skip lists). The informal and intuitive probabilistic analysis we give below uses only basic concepts of probability theory. Let us begin by determining the expected value of the height of a skip list S with n entries (assuming that we do not terminate insertions early). The probability that a given entry has a tower of height $i+1$ is equal to the probability of getting i consecutive heads when flipping a coin, that is, this probability is $1/2^i$.

Hence, the

probability that level i has at least one position is at most

$$1/2^{i-1}$$

for

the probability that any one of n different events occurs is at most the sum of the probabilities that each occurs.

The probability that the height of S is larger than i is equal to the probability that level i has at least one position, that is, it is no more than $1/2^{i-1}$. This means that S is larger than, say, $3 \log n$ with probability at most

$$1/2^{3 \log n - 1}$$

$$1/n^2$$

$$= 1/n^2$$

$$1/n^3 = 1/n^3$$

$$1/n^2$$

For example, if $n=1000$, this probability is a one-in-a-million long shot. More generally, given a constant $c > 1$, S is larger than $c \log n$ with probability at most $1/n^{c-1}$. That is, the probability that S is smaller than $c \log n$ is at least $1 - 1/n^{c-1}$.

Thus, with high probability, the height of S is $O(\log n)$.

Analyzing Search Time in a Skip List

Next, consider the running time of a search in skip list S , and recall that such a search involves two nested while loops. The inner loop performs a scan forward on all levels of S as long as the next key is no greater than the search key k , and the outer loop drops down to the next level and repeats the scan forward iteration. Since the height of S is $O(\log n)$ with high probability, the search time is $O(\log n)$ with high probability.

10.4. Skip Lists 443

Maintaining the Topmost Level

A skip list must maintain a reference to the start position (the topmost, left position in S) as an instance variable, and must have a policy for any insertion that wishes to continue inserting a new entry past the top level of S . There are two possible courses of action we can take, both of which have their merits.

One possibility is to restrict the top level, h , to be kept at some fixed value that is a function of n , the number of entries currently in the map (from the analysis we will see that $h = \max\{10, 2 \cdot \log n\}$ is a reasonable choice, and picking $h = 3 \cdot \log n$ is even safer). Implementing this choice means that we must modify the insertion algorithm to stop inserting a new position once we reach the topmost level (unless $\log n < \log(n+1)$, in which case we can now go at least one more level, since the bound on the height is increasing).

The other possibility is to let an insertion continue inserting a new position as long as heads keeps getting returned from the random number generator. This is the approach taken by algorithm in the analysis of skip lists, the probability that an insertion will go to a level that is more than $O(\log n)$ is very low, so. Either choice will still result in the expected $O(\log n)$ time to perform search, insertion, and removal, however, which we show in the next section.

10.4.2 Probabilistic Analysis of Skip Lists

As we have shown above, skip lists provide a simple implementation of a sorted map. In terms of worst-case performance, the probability of having a fair coin repeatedly come up heads forever is 0. Moreover, we cannot infinitely add positions to a list without eventually running out of memory. In any case, if we terminate position, the worst-case running time for performing the

getitem

,

setitem

, and

delitem

map operations in a skip list with n entries and height h is $O(n+h)$. This worst-case performance occurs when the tower of every entry reaches level $h-1$, where h is the height of S . However, this event has very low probability. Judging from this worst case, we might conclude that the skip-list structure is strictly inferior to the other map implementations. But, as we have seen, this worst-case behavior is a gross overestimate.

10.4. Skip Lists 439

10.4.1 Search and Update Operations in a Skip List

The skip-list structure affords simple map search and update algorithms. In fact, all of the skip-list search and update algorithms are based on an elegant SkipSearch method that takes a key k and finds the position p of the item in list S that has the largest key less than or equal to k (which is possibly \perp).

Searching in a Skip List

Suppose we are given a search key k . We begin the SkipSearch method by setting a position variable p to the topmost, left position in the skip list S , called the start position of S . That is, the start position is the position of S storing the special entry with key \perp . We then perform the following steps (see Figure 10.11), where $\text{key}(p)$ denotes the key of the item at position p :

1. If $S.\text{below}(p)$ is None , then the search terminates—we are at the bottom and have located the item in S with the largest key less than or equal to the search key k . Otherwise, we drop down to the next lower level in the present tower by setting $p = S.\text{below}(p)$.

2. Starting at position p , we move p forward until it is at the rightmost position on the present level such that $\text{key}(p) \leq k$. We call this the scan forward step.

Note that such a position always exists, since each level contains the keys \perp and ∞ . It may be that p remains where it was after the scan forward for this level.

3. Return to step 1.

55 S1 S2 S3 S4 S5

+•

+•+•

+•

+•+•

---•

-- 12 12 -- 17

17 25

25 20 17 31 38 39 --

-- 17 17 25

25 31

31 38 44

44 50 55 55 55

S0

Figure 10.11: Example of a search in a skip list. The positions examined when searching for key 50 are highlighted.

We give a pseudo-code description of the skip-list search algorithm, SkipSearch, in Code Fragment 10.12. Given this method, the map operation $M[k]$ is performed by computing $p = \text{SkipSearch}(k)$ and testing whether or not $\text{key}(p) = k$. If these two keys are equal, we return the associated value; otherwise, we raise a `KeyError`.

10.4. Skip Lists 441

Algorithm SkipInsert(k,v) :

Input: Key and value v

Output: Topmost position of the item inserted in the skip list

p=SkipSearch (k)

q=None {q will represent top node in new item's tower }

i=1

repeat

i=i+1

if i=h then

h=h+1 {add a new level to the skip list }

t=next(s)

s=insertAfterAbove (None,s,(.,None)) {grow leftmost tower }

insertAfterAbove (s,t,(.,None)) {grow rightmost tower }

while above (p) is None do

p=prev (p) {scan backward }

p=above (p) {jump up to higher level }

q=insertAfterAbove (p,q,(k,v)){increase height of new item's tower }

until coinFlip ()==tails

n=n+1

return q

Code Fragment 10.13: Insertion in a skip list. Method coinFlip () returns ·heads· or ·tails·, each with probability 1 /2. Instance variables n,h,a n d s hold the number of entries, the height, and the start node of the skip list.

55 S1S2S3S4S5

+·

+·+·

+·

+·+·

-- 1212 --17

17 25

25 20 17 31--

-- 1717 25

25 31

31 38 44

44424242

5555

55 38 39 42 50 S0

Figure 10.12: Insertion of an entry with key 42 into the skip list of Figure 10.10. We assume that the random ·coin ·ips· for the new entry came up heads three times in a row, followed by tails. The positions visited are highlighted. The positions inserted to hold the new entry are drawn

718 Chapter 15. Memory Management and B-Trees

C-15.11 Describe an external-memory data structure to implement the queue ADT so that the total number of disk transfers needed to process a sequence of enqueue and dequeue operations is $O(k/B)$.

C-15.12 Describe an external-memory version of the PositionalList ADT (Section 7.4), with block size B , such that an iteration of a list of length n is completed using $O(n/B)$ transfers in the worst case, and all other methods of the ADT require only $O(1)$ transfers.

C-15.13 Change the rules that define red-black trees so that each red-black tree T has a corresponding $(4,8)$ tree, and vice versa.

C-15.14 Describe a modified version of the B-tree insertion algorithm so that each time we create an overflow block, we split the keys among all of w 's siblings, so that each sibling holds roughly the same number of keys (possibly cascading the split up to the parent of w). What is the minimum fraction of each block that will always be filled using this scheme?

C-15.15 Another possible external-memory map implementation is to use a skip list, but to collect consecutive groups of $O(B)$ nodes, in individual blocks, on any level in the skip list. In particular, we define an order- d B-skip list to be such a representation of a skip list structure, where each block contains at least $d/2$ list nodes and at most d list nodes. Let us also choose d in this case to be the maximum number of list nodes from a level of a skip list that can fit into one block. Describe how we should modify the skip-list insertion and removal algorithms for a B-skip list so that the expected height of the structure is $O(\log n / \log B)$.

C-15.16 Describe how to use a B-tree to implement the partition (union-find) ADT (from Section 14.7.3) so that the union and find operations each use at most $O(\log n / \log B)$ disk transfers.

C-15.17 Suppose we are given a sequence S of n elements with integer keys such that some elements in S are colored blue and some elements in S are colored red. In addition, say that a red element e pairs with a blue element f if they have the same key value. Describe an efficient external-memory algorithm for finding all the red-blue pairs in S . How many disk transfers does your algorithm perform?

C-15.18 Consider the page caching problem where the memory cache can hold m pages, and we are given a sequence P of n requests taken from a pool of $m+1$ possible pages. Describe the optimal strategy for the offline algorithm and show that it causes at most $m + n/m$ page misses in total, starting from an empty cache.

C-15.19 Describe an efficient external-memory algorithm that determines whether an array of n integers contains a

10.4. Skip Lists 445

So we have yet to bound the number of scan-forward steps we make. Let n_i be the number of keys examined while scanning forward at level i . Observe that, after the key at the starting position, each additional key examined in a scan-forward at level i cannot also belong to level $i+1$. If any of these keys were on the previous level, we would have encountered them in the previous scan-forward step. Thus, the probability that any key is co-
 n_i is $1/2$. Therefore, the expected value of n_i is exactly equal to the expected number of times we must flip a fair coin before it comes up heads. This expected value is 2. Hence, the expected amount of time spent scanning forward at any probability, a search in Stakes expected time $O(\log n)$. By a similar analysis, we can show that the expected running time of an insertion or a removal is $O(\log n)$.

Space Usage in a Skip List

Finally, let us turn to the space requirement of a skip list with n entries. As we observed above, the expected number of positions at level i is $n/2^i$, which means that the expected total number of positions in S is

$$\sum_{i=0}^h \frac{n}{2^i}$$

Using Proposition 3.5 on geometric summations, we have

$$\sum_{i=0}^h \frac{n}{2^i} = \frac{n}{2^0} + \frac{n}{2^1} + \frac{n}{2^2} + \dots + \frac{n}{2^h} < 2n$$

Hence, the expected space requirement of S is $O(n)$.

Table 10.4 summarizes the performance of a sorted map realized by a skip list.

Operation

Running Time

$\text{len}(M)$

$O(1)$

$\text{kin}M$

$O(\log n)$ expected

$M[k] = v$

$O(\log n)$ expected

$\text{del } M[k]$

$O(\log n)$ expected

$M.\text{nd}$

$\text{min}(), M.\text{nd}$

$\text{max}()$

$O(1)$

$M.\text{nd}$